

Chapter 3: Processes



Chapter 3: Processes

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Inter-Process Communication (IPC)
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems



Objectives

- ❑ Identify the separate **components of a process** and illustrate how they are represented and scheduled in an operating system.
- ❑ Describe **how processes are created and terminated** in an operating system, including developing programs using the appropriate system calls that perform these operations.
- ❑ Describe and contrast **inter-process communication** using shared memory and message passing.
- ❑ Design *programs that uses pipes and POSIX shared memory* to perform inter-process communication.
- ❑ Describe **client-server communication** using sockets and remote procedure calls.
- ❑ Design *kernel modules* that interact with the Linux operating system.



Process Concept

- ❑ An operating system executes a variety of programs that run as processes
- ❑ **Process** – a program in execution; process execution must progress in sequential fashion
- ❑ Multiple parts
 - The *program code*, also called *text section*
 - Current activity including *program counter*, and *processor registers*
 - *Stack section* containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - *Data section* containing global variables
 - *Heap section* containing memory dynamically allocated during run time




Process Concept (Cont.)

- ❑ *Program* is *passive* entity stored on disk (e.g., *executable file*)
- ❑ *Process* is *active* entity
 - Program becomes process when executable file loaded into memory
- ❑ *Execution of program* can be started via GUI mouse clicks, command line (CLI) entry of its name, etc.
- ❑ One program can be several processes
 - E.g., Consider multiple users executing the same program

`#ps -aux`



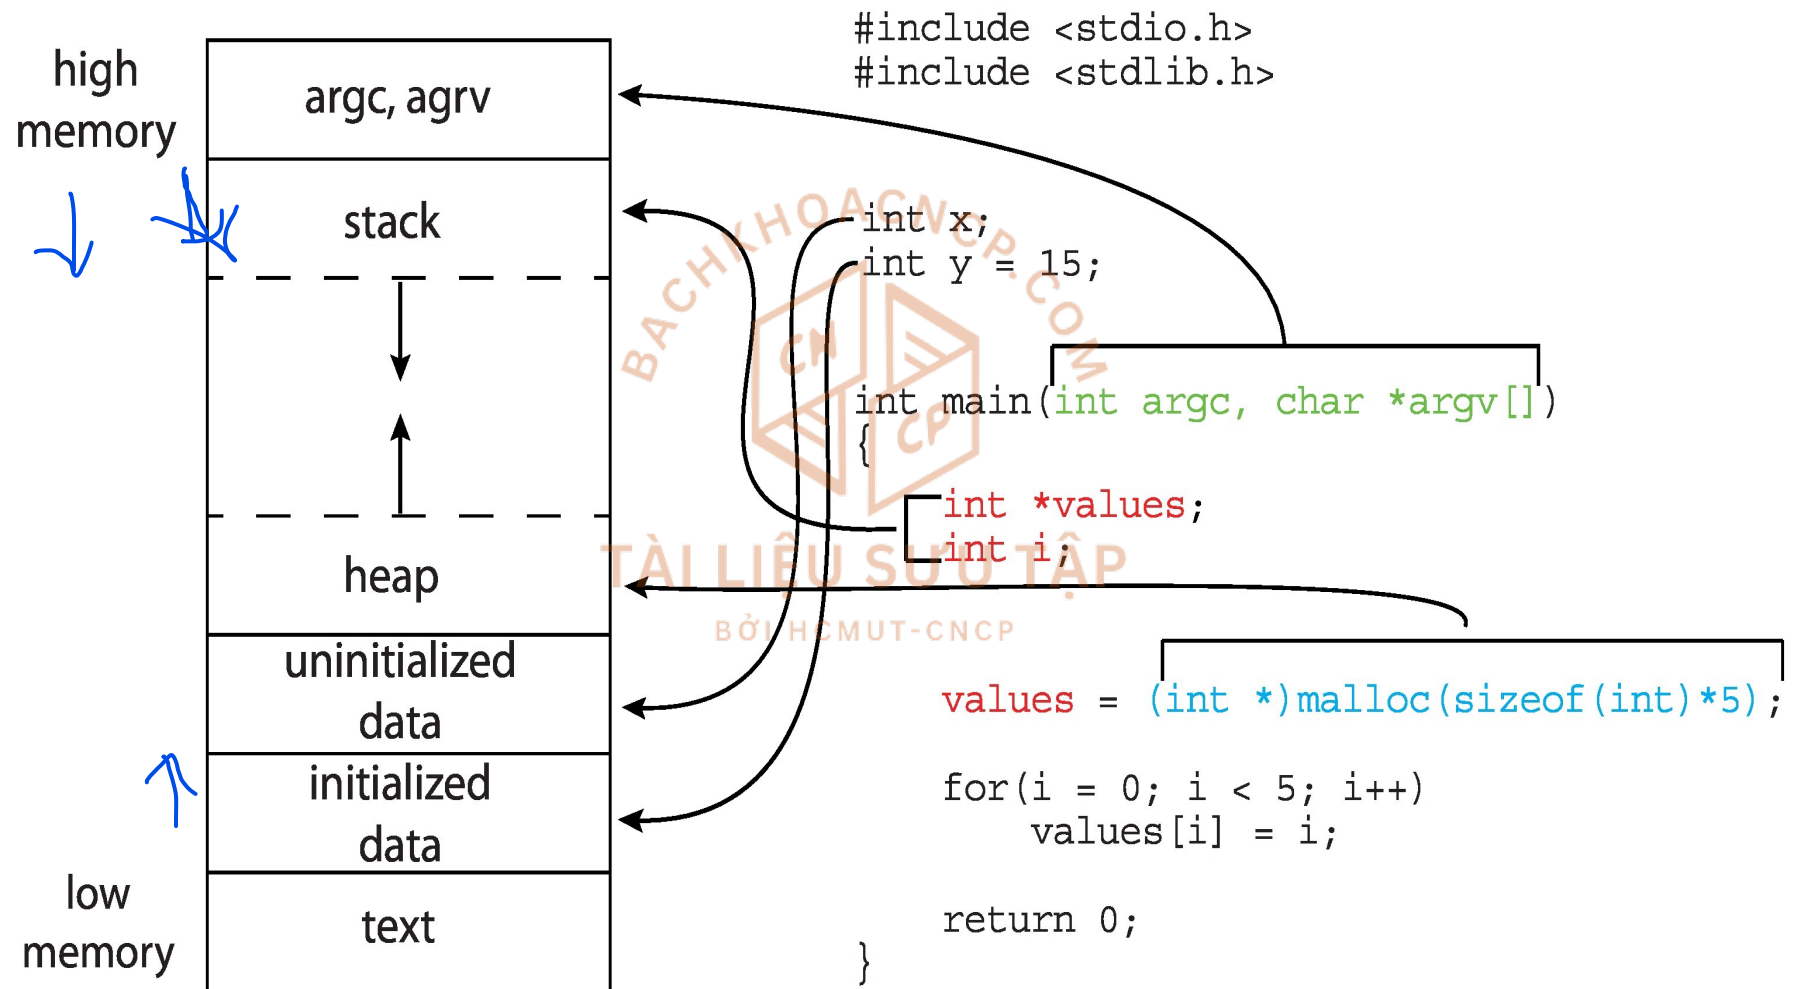
Process in Memory



`#size <pid>`



Memory Layout of a C Program

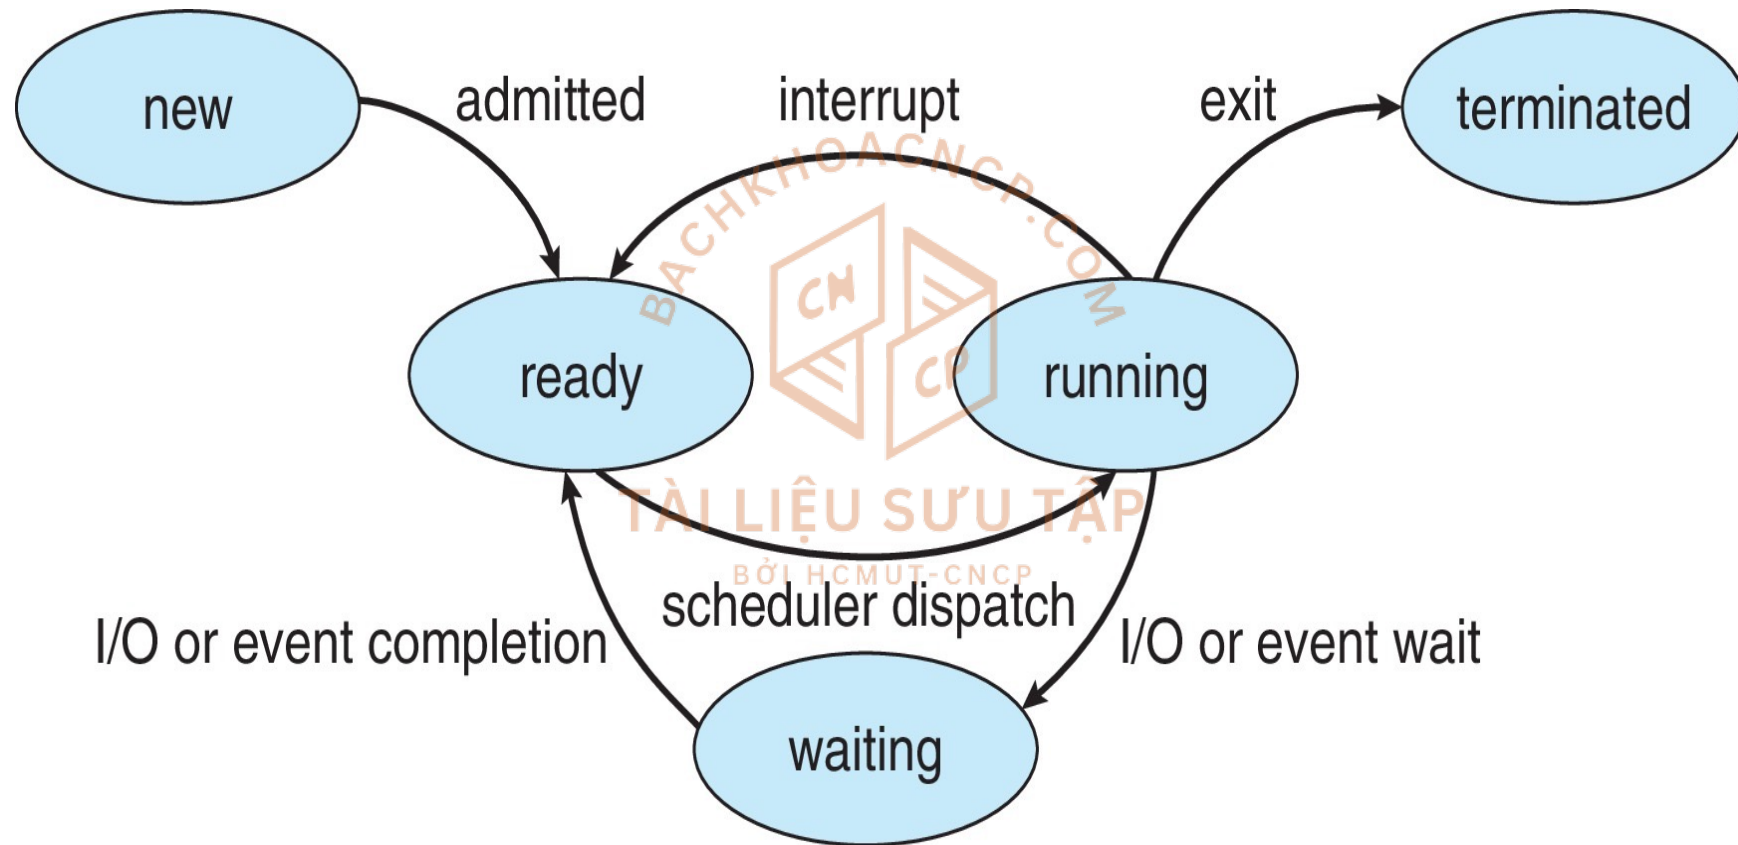


Process State

- ❑ As a process executes, it changes *state*
 - *New* – The process is being created
 - *Running* – Instructions are being executed
 - *Waiting* – The process is waiting for some event to occur
 - *Ready* – The process is waiting to be assigned to a processor
 - *Terminated* – The process has finished execution



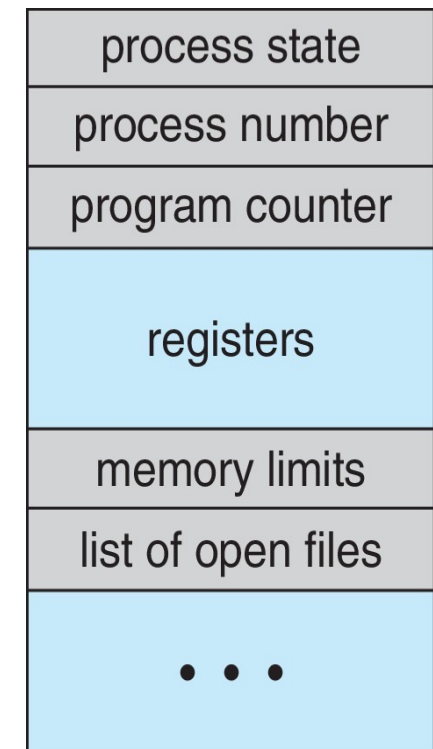
Diagram of Process State



Process Control Block (PCB)

❑ **Process Control Block (PCB)** – Information associated with each process, also called **Task Control Block (TCB)**, includes:

- *Process state* – running, waiting, etc.
- *Process number* – identity of the process
- *Program counter* – location of instruction to next execute
- *CPU registers* – contents of all process-centric registers
- *CPU scheduling info* – priorities, scheduling queue pointers
- *Memory-management information* – memory allocated to the process
- *Accounting information* – CPU used, clock time elapsed since start, time limits
- *I/O status information* – I/O devices allocated to process, list of open files



Threads

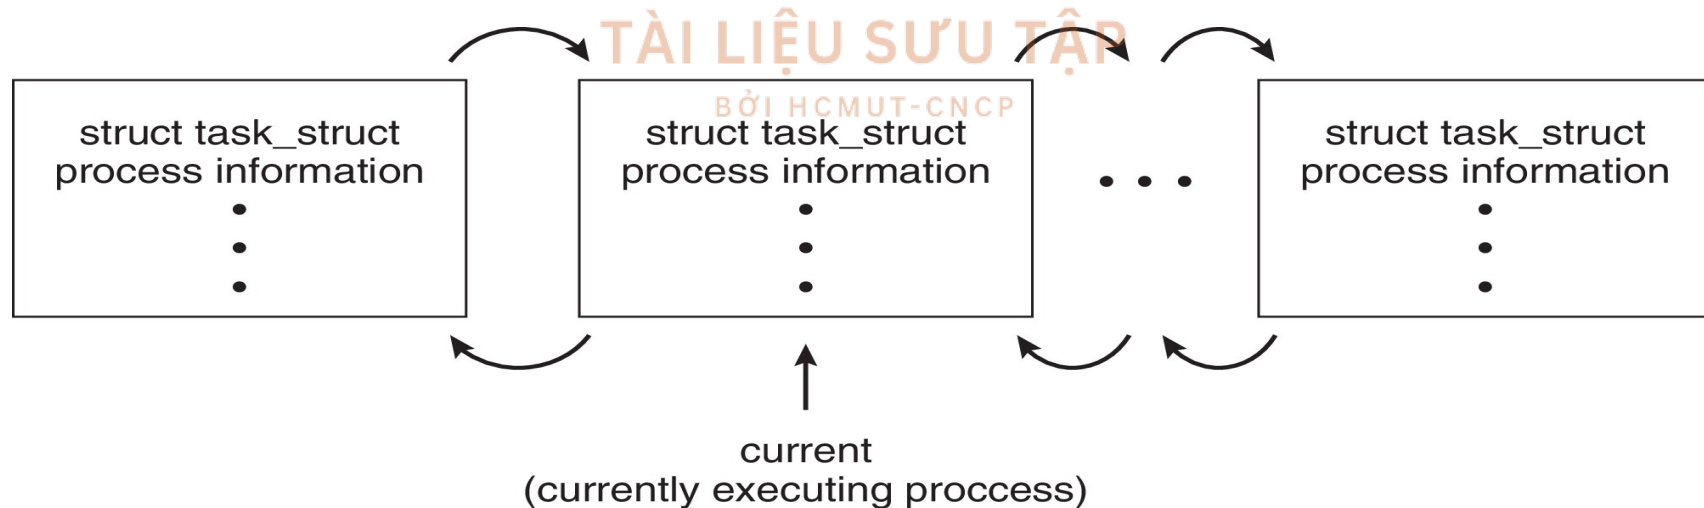
- ❑ So far, process has a *single thread* of execution
- ❑ Consider having *multiple program counters per process*
 - Multiple locations can execute at once
 - ▶ Multiple threads of control → *threads*
- ❑ Must then have *storage for thread* details
- ❑ Multiple program counters in PCB



Process Representation in Linux

□ Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```




Process Scheduling

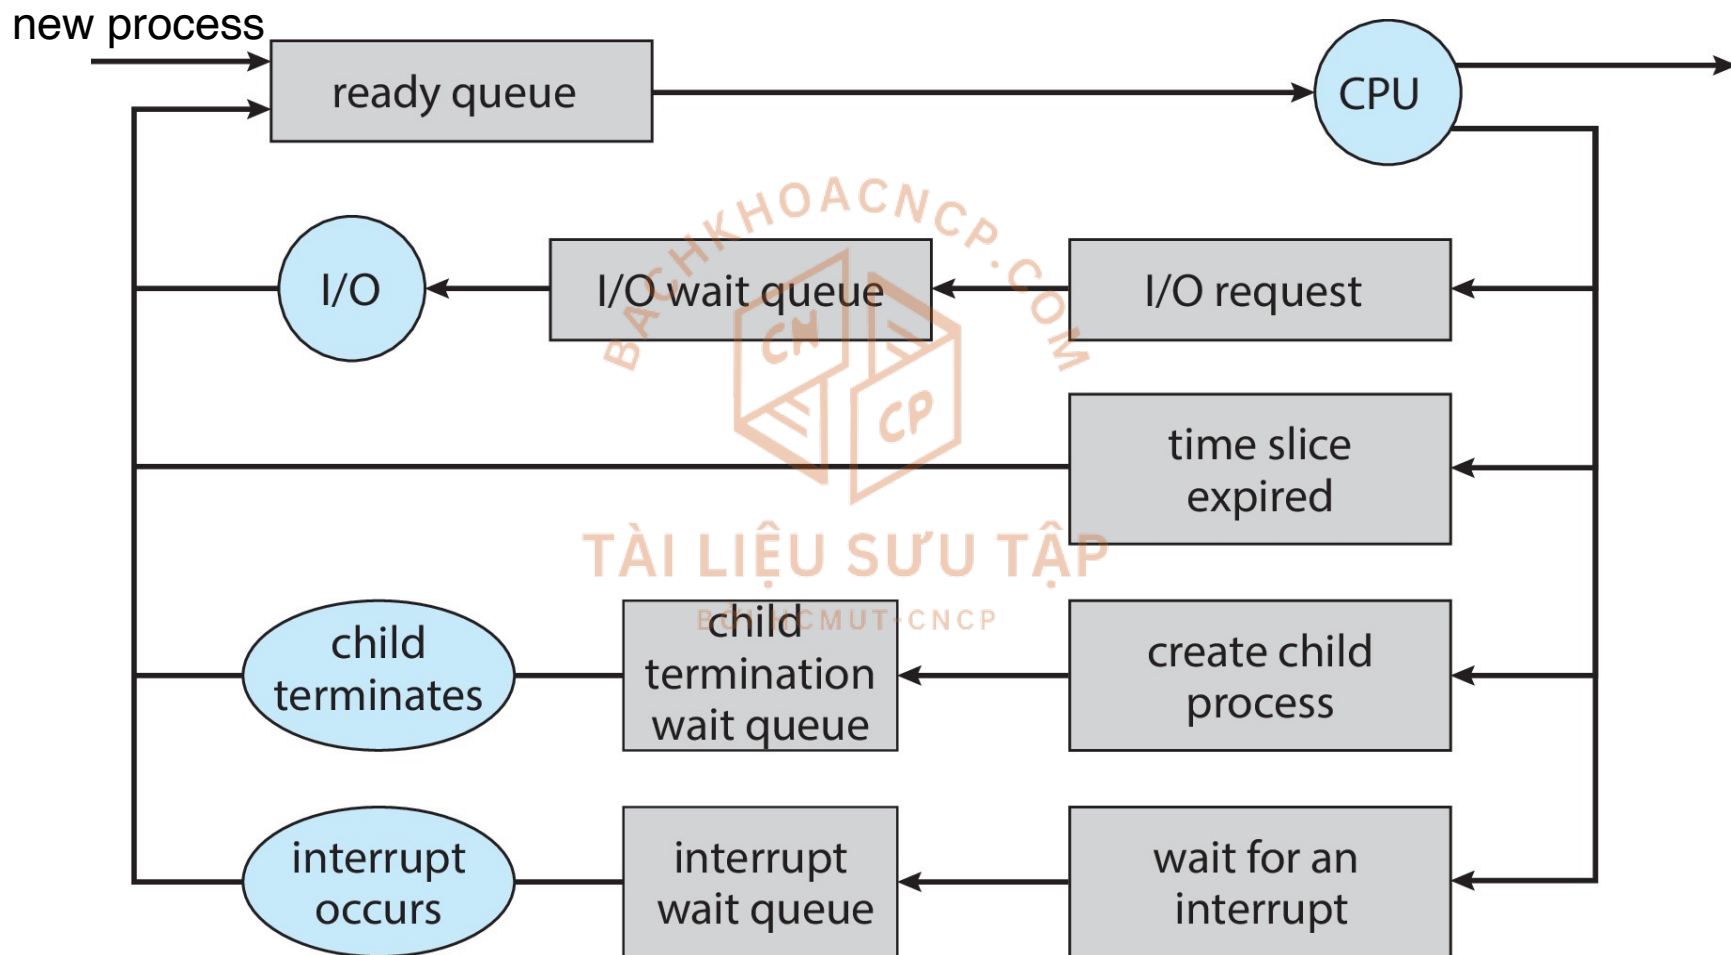
- ❑ Maximize CPU use → quickly switch processes onto CPU core
- ❑ *Process scheduler* selects one process among available (ready) processes for next execution on CPU core
- ❑ Maintains *scheduling queues* of processes
 - *Ready queue* – set of all processes residing in main memory, ready and waiting to execute
 - *Wait queues* – set of processes waiting for an event (e.g., I/O)
- ❑ Processes migrate among the various queues



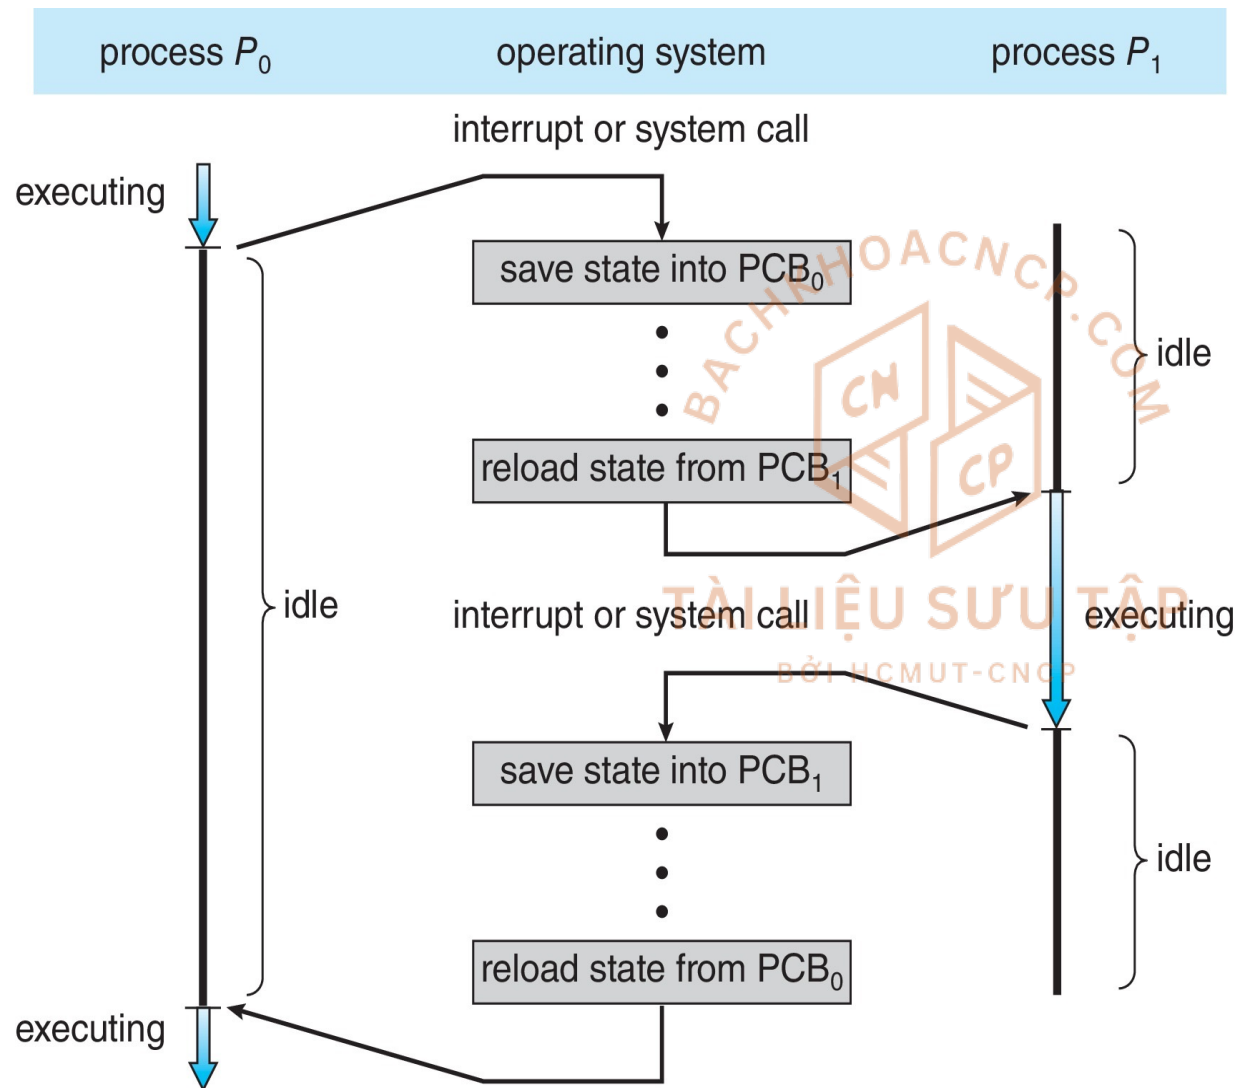
Ready and Wait Queues



Representation of Process Scheduling



CPU Switch from Process to Process



- A *context switch* occurs when the CPU switches from one process to another.



Context Switch

- ❑ When CPU switches to another process, the system must *save the state* of the old process and load the *saved state* for the new process via a *context switch*
- ❑ *Context* of a process represented in the **PCB**
- ❑ Context-switch time is *overhead*, the system does no useful work while switching
 - The more complex the OS and the PCB, the longer the context switch
- ❑ *Time* dependent on hardware support
 - Some hardware provides *multiple sets of registers per CPU*, multiple contexts loaded at once



Operations on Processes

❑ System must provide mechanisms for:

- process creation
- process termination



Process Creation

- ❑ *Parent processes* create *children processes*, which, in turn create other processes, forming a *tree of processes*
- ❑ Process identified and managed via a **Process Identifier (PID)**
- ❑ **Resource sharing options**
 - Parent and children share *all* resources
 - Children share *subset* of parent's resources
 - Parent and child share *no* resources




Process Creation (Cont.)

❑ Execution options

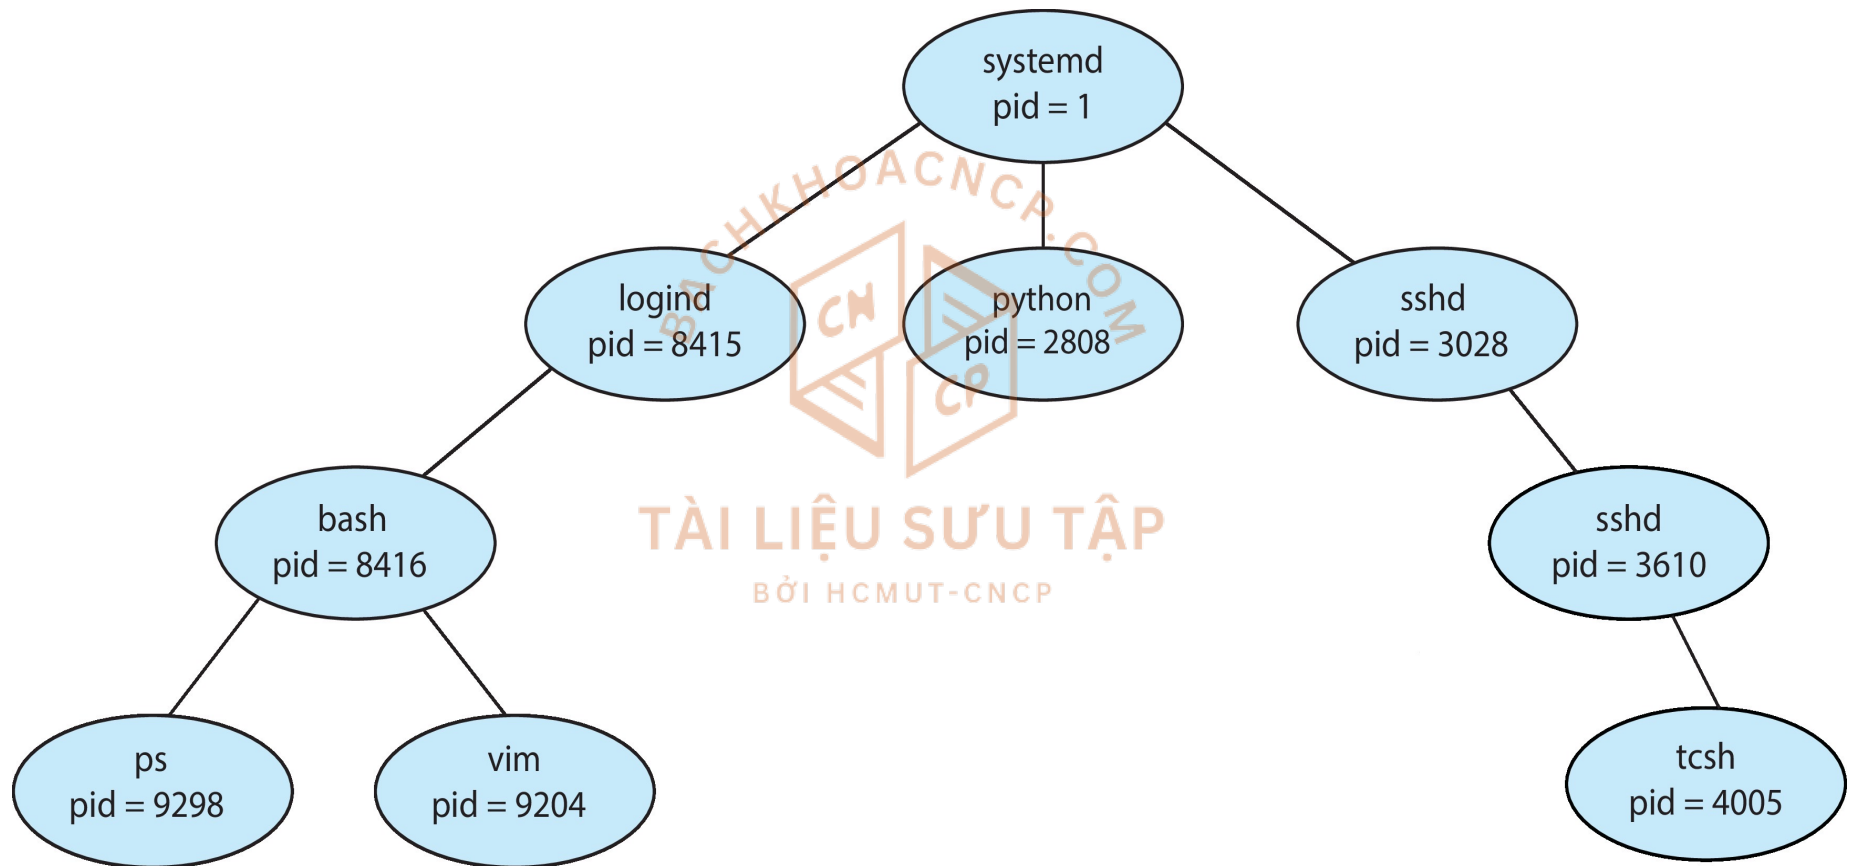
- Parent and children execute concurrently
- Parent waits until children terminate

❑ Address space

- Child duplicate of parent
- Child has a program loaded into it



A Tree of Processes in Linux



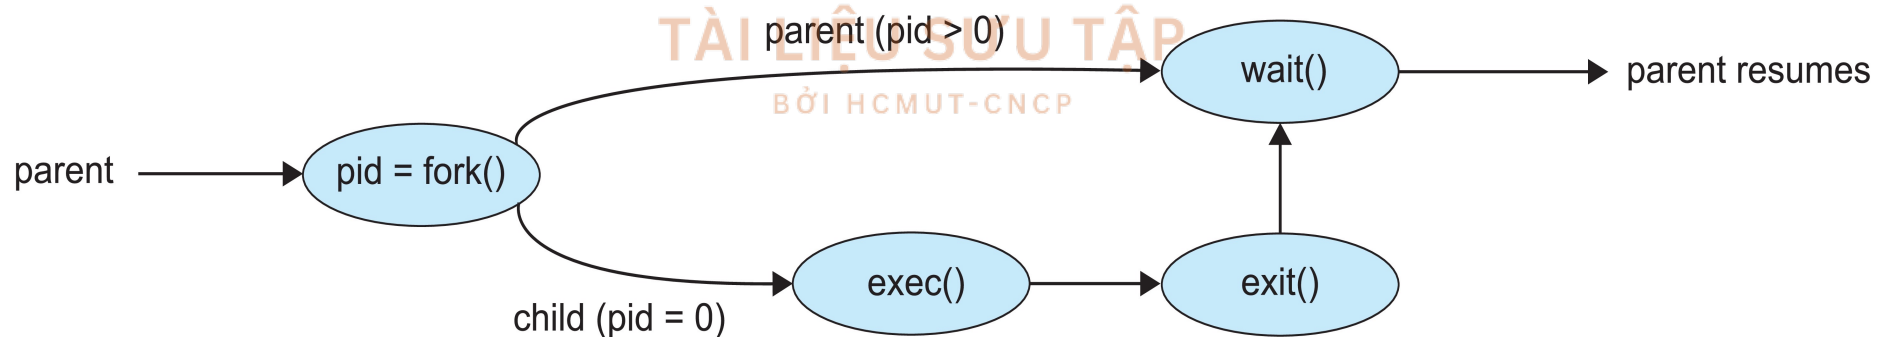
#pstree



Process Creation (Cont.)

❑ UNIX examples

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program
- Parent process calls **wait()** waiting for the child to terminate



C Program Forking A Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```



Example

❑

How many processes will be generated?

❑ int main(){

❑ printf("Hello \n");

❑ 1. fork();

❑ 2. fork();

❑ printf("Hello \n");

❑ return 0;

❑ }



Process Termination

- ❑ Process executes *last statement* and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- ❑ Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the *operating systems does not allow a child to continue if its parent terminates*




Process Termination (Cont.)

- ❑ Some operating systems do not allow child to exist if its parent has terminated. *If a process terminates, then all its children must also be terminated.*
 - **Cascading termination:** All children, grandchildren, etc. are terminated
 - The termination is initiated by the operating system
 - ❑ The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the **pid** of the terminated process
- ```
pid = wait(&status);
```
- ❑ If no parent waiting (did not invoke `wait()`), process is a *zombie*
  - ❑ If parent terminated without invoking `wait()`, process is an *orphan*



# Multiprocess Architecture – Chrome Browser

- ❑ Many web browsers ran as a single process (some still do)
  - If one web site causes trouble, entire browser can hang or crash
- ❑ Google Chrome Browser is multiprocess with 3 different types of processes:
  - *Browser process* manages user interface, disk and network I/O
  - *Renderer process* renders web pages, deals with HTML, JavaScript. A new renderer created for each website opened
    - ▶ Runs in *sandbox* restricting disk and network I/O, minimizing effect of security exploits
  - *Plug-in process* for each type of plug-in



Each tab represents a separate process.



# Inter-Process Communication (IPC)

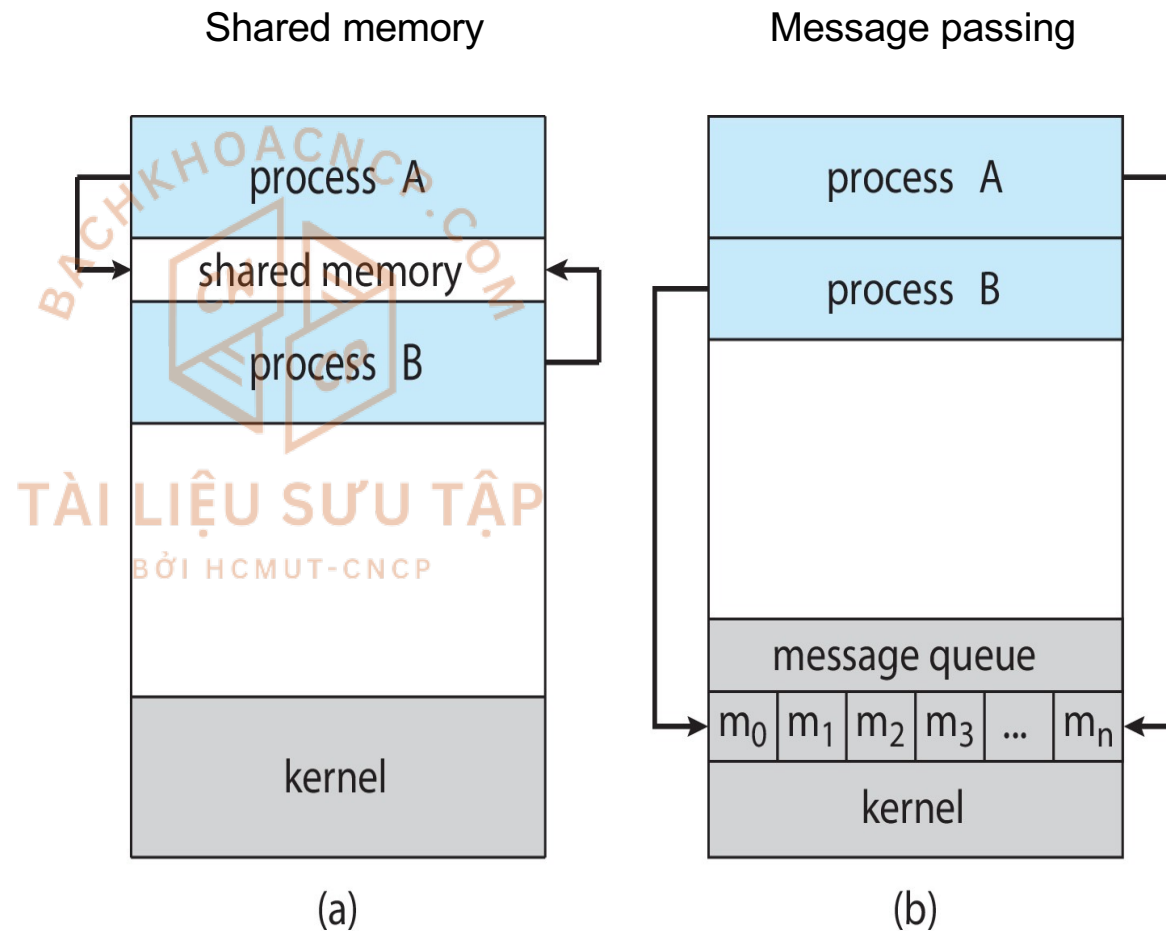
- ❑ Processes within a system may be *independent* or *cooperating*
  - *Independent process* does not share data with any other processes executing in the system
  - *Cooperating process* can affect or be affected by other processes, including sharing data
- ❑ Reasons for cooperating processes:
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience
- ❑ Cooperating processes need **Inter-Process communication (IPC)**



# Communication Models

## □ Two models of IPC

- *Shared memory*
- *Message passing*



# Inter-Process Communication – Shared Memory

- ❑ An *area of memory shared among the processes* that wish to communicate
- ❑ The communication is *under the control of the users processes*, not the operating system.
- ❑ Major issues is to provide mechanism that will allow the user processes to *synchronize their actions* when they access shared memory.

TÀI LIỆU SƯU TẬP  
BỞI HCMUT-CNCP



# Producer-Consumer Problem

- ❑ *Producer-Consumer relationship*
- ❑ Paradigm for cooperating processes, *producer process* produces information that is consumed by a *consumer process*
  - *unbounded-buffer* places no practical limit on the size of the buffer
  - *bounded-buffer* assumes that there is a fixed buffer size

TÀI LIỆU SƯU TẬP  
BỞI HCMUT-CNCP



# Bounded-Buffer – Shared-Memory Solution

## ❑ Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
 . . .
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

## ❑ Solution is correct, but can only use **BUFFER\_SIZE-1** elements





# Producer Process – Shared Memory

```
item next_produced;

while (true) {
 /* produce an item in next produced */
 while (((in + 1) % BUFFER_SIZE) == out)
 ; /* do nothing */

 buffer[in] = next_produced;

 in = (in + 1) % BUFFER_SIZE;
}
```



# Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
 while (in == out)
 ; /* do nothing */
 next_consumed = buffer[out];
 out = (out + 1) % BUFFER_SIZE;

 /* consume the item in next consumed */
}
```



# Inter-Process Communication – Message Passing

---

- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
  - send(message)
  - receive(message)
- ❑ The message size is either fixed or variable



# Message Passing (Cont.)

- ❑ If processes P and Q wish to communicate, they need to:
  - Establish a communication link between them
  - Exchange messages via **send/receive**
- ❑ Implementation issues:
  - How are links established?
  - Can a link be associated with more than two processes?
  - How many links can there be between every pair of communicating processes?
  - What is the capacity of a link?
  - Is the size of a message that the link can accommodate fixed or variable?
  - Is a link unidirectional or bi-directional?



# Direct Communication

- ❑ Processes must name each other explicitly:
  - `send(P, message)` – send a message to process P
  - `receive(Q, message)` – receive a message from process Q
- ❑ Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional



# Indirect Communication

- ❑ Messages are directed and received from *mailboxes* (also referred to as *ports*)
  - Each mailbox has a *unique ID*
  - Processes can communicate *only if they share a mailbox*
- ❑ Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional



# Indirect Communication (Cont.)

## □ Operations

- create a new mailbox (or port)
- send and receive messages through mailbox
- destroy a mailbox

## □ **Primitives** are defined as:

- **send**(*A, message*) – send a message to mailbox A
- **receive**(*A, message*) – receive a message from mailbox A



# Indirect Communication (Cont.)

## □ Mailbox sharing

### ○ Example

- ▶  $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A,
- ▶  $P_1$  sends;  $P_2$  and  $P_3$  receive.
- ▶ Who gets the message?

### ○ Solutions

- ▶ Allow a link to be associated with at most two processes
- ▶ Allow only one process at a time to execute a receive operation
- ▶ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was





# Message Passing – Synchronization

- ❑ Message passing may be either *blocking* or *non-blocking*
- ❑ *Blocking* is considered *synchronous*
  - *Blocking send* – the sender is blocked until the message is received
  - *Blocking receive* – the receiver is blocked until a message is available
- ❑ *Non-blocking* is considered *asynchronous*
  - *Non-blocking send* – the sender sends the message and continue
  - *Non-blocking receive* – the receiver receives:
    - ▶ A valid message, or Null message
- ❑ Different combinations possible
  - If both send and receive are blocking, we have a *rendezvous*



# Producer – Message Passing

```
message next_produced;

while (true) {
 /* produce an item in next_produced */

 send(next_produced);
}
```



# Consumer – Message Passing

```
message next_consumed;

while (true) {
 receive(next_consumed)

 /* consume the item in next_consumed */
}
```



# Buffering

- ❑ Queue of messages attached to the link.
- ❑ Implemented in one of three ways
  - *Zero capacity* – no messages are queued on a link
    - ▶ Sender must wait for receiver (rendezvous)
  - *Bounded capacity* – finite length of  $n$  messages
    - ▶ Sender must wait if link full
  - *Unbounded capacity* – infinite length
    - ▶ Sender never waits



# Examples of IPC Systems - POSIX

## ❑ POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment
- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Use `mmap()` to memory-map a file pointer to the shared memory object
- Reading and writing to shared memory is done by using the pointer returned by `mmap()`.




# Pipes

- ❑ Acts as a conduit allowing two processes to communicate
- ❑ Issues:
  - Is communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half or full-duplex?
  - Must there exist a relationship (e.g., *parent-child*) between the communicating processes?
  - Can the pipes be used over a network?
- ❑ *Ordinary pipes* – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- ❑ *Named pipes* – can be accessed without a parent-child relationship.



# Ordinary Pipes

- ❑ *Ordinary Pipes* allow communication in standard producer-consumer style
  - *Producer* writes to one end (the *write-end* of the pipe)
  - *Consumer* reads from the other end (the *read-end* of the pipe)
- ❑ Ordinary pipes are therefore unidirectional
- ❑ Require *parent-child relationship* between communicating processes



# Named Pipes

---

- ❑ *Named pipes* are more powerful than ordinary pipes
- ❑ Communication is bidirectional
- ❑ *No parent-child relationship* is necessary between the communicating processes
- ❑ Several processes can use the named pipe for communication
- ❑ Provided on both **UNIX** and **Windows** systems





# Communications in Client-Server Systems

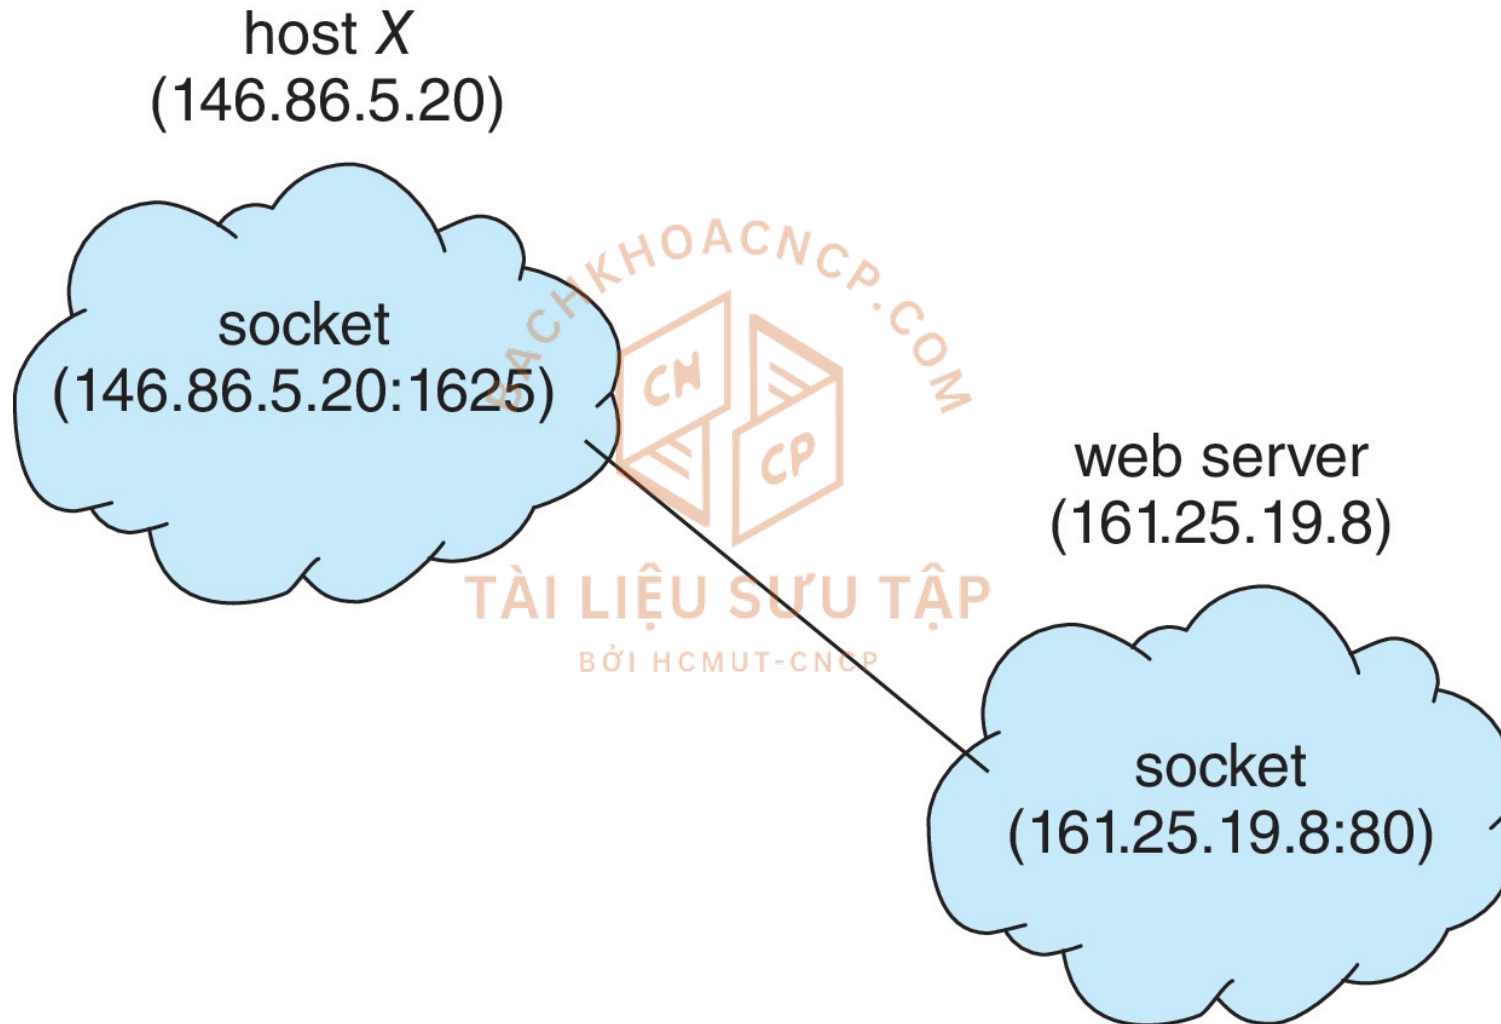
## ❑ Sockets

- A *socket* is defined as an endpoint for communication
- It is a concatenation of *IP address* and *port* – a number included at start of message packet to differentiate network services on a host
  - ▶ E.g., The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are *well known*, used for standard services
- Special IP address **127.0.0.1** (*loopback*) to refer to system on which process is running

## ❑ Remote Procedure Calls (RPC)



# Socket Communication



# Sockets in Java – Server

```
import java.net.*;
import java.io.*;

public class DateServer
{
 public static void main(String[] args) {
 try {
 ServerSocket sock = new ServerSocket(6013);

 /* now listen for connections */
 while (true) {
 Socket client = sock.accept();

 PrintWriter pout = new
 PrintWriter(client.getOutputStream(), true);

 /* write the Date to the socket */
 pout.println(new java.util.Date().toString());

 /* close the socket and resume */
 /* listening for connections */
 client.close();
 }
 } catch (IOException ioe) {
 System.err.println(ioe);
 }
 }
}
```

## □ Three types of sockets

- *Connection-oriented (TCP)*
- *Connectionless (UDP)*
- **MulticastSocket** class— data can be sent to multiple recipients

## □ Consider this “Date” *server* in Java:



# Sockets in Java – Client

```
import java.net.*;
import java.io.*;

public class DateClient
{
 public static void main(String[] args) {
 try {
 /* make connection to server socket */
 Socket sock = new Socket("127.0.0.1",6013);

 InputStream in = sock.getInputStream();
 BufferedReader bin = new
 BufferedReader(new InputStreamReader(in));

 /* read the date from the socket */
 String line;
 while ((line = bin.readLine()) != null)
 System.out.println(line);

 /* close the socket connection*/
 sock.close();
 }
 catch (IOException ioe) {
 System.err.println(ioe);
 }
 }
}
```

- The equivalent  
“Date” *client*



# Remote Procedure Calls

- ❑ **Remote Procedure Call (RPC)** abstracts procedure calls between processes on *networked systems*
  - Again uses *ports* for service differentiation
- ❑ **Stubs** – proxies for the actual procedure on the server and client sides
  - The *client-side stub* locates the server and *marshals* the parameters
  - The *server-side stub* receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- ❑ On **Windows**, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**



# Remote Procedure Calls (Cont.)

- ❑ Data representation handled via **External Data Representation (XDR)** format to account for different architectures
  - E.g., *Big-endian* (Motorola) and *little-endian* (Intel x86)
- ❑ Remote communication has *more failure scenarios* than local
  - Messages can be delivered *exactly once* rather than *at most once*
- ❑ OS typically provides a *rendezvous* (or *matchmaker*) service to connect client and server



# End of Chapter 3

