

Chapter 4: Threads & Concurrency



Chapter 4: Threads

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples



Objectives

- ❑ Identify the *basic components of a thread*, and *contrast threads and processes*
- ❑ Describe the *benefits* and *challenges* of designing *multithreaded applications*
- ❑ Illustrate *different approaches to implicit threading* including thread pools, fork-join, and Grand Central Dispatch
- ❑ Describe how the Windows and Linux operating systems represent threads
- ❑ Design multithreaded applications using the Pthreads, Java, and Windows threading APIs

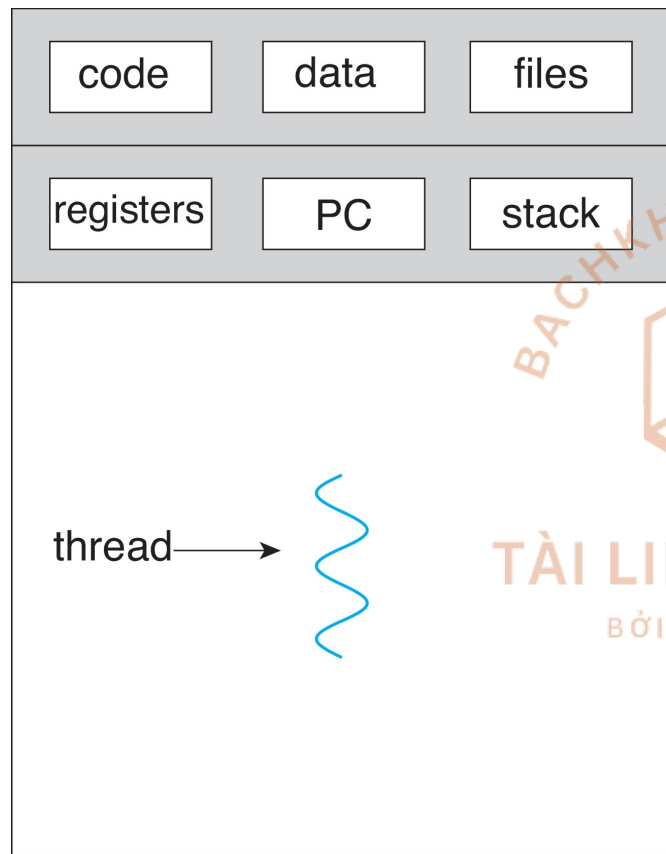


Motivation

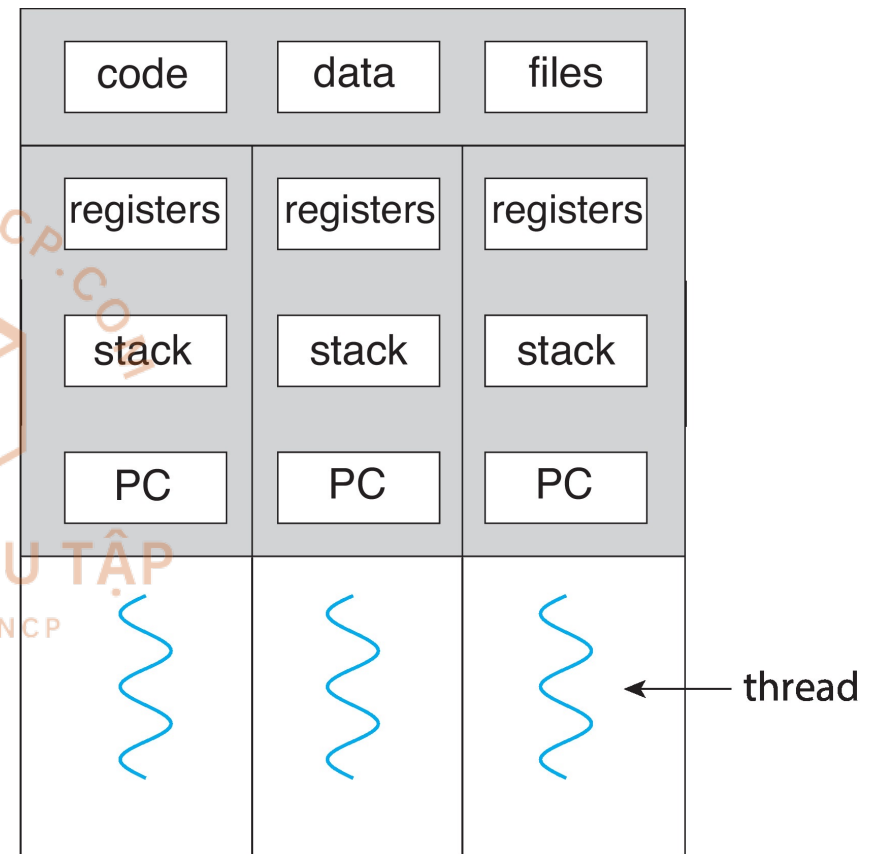
- ❑ Most modern applications are *multithreaded*
- ❑ *Threads run within application*
- ❑ Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- ❑ Process creation is heavy-weight while *thread creation is light-weight*
- ❑ Can simplify code, increase efficiency
- ❑ Kernels are generally multithreaded



Single and Multithreaded Processes



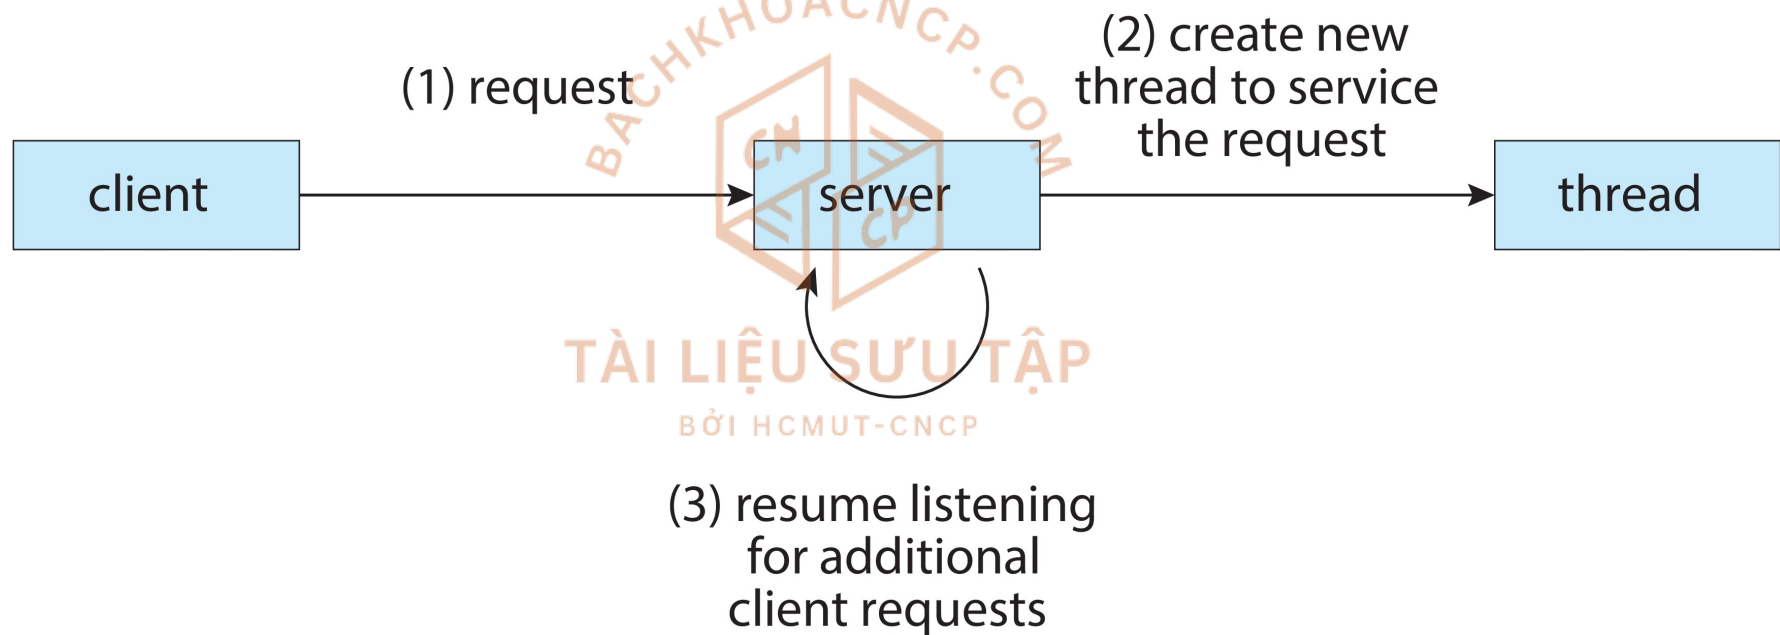
single-threaded process



multithreaded process



Multithreaded Server Architecture



Benefits

- ❑ *Responsiveness* – may allow continued execution if part of process is blocked, especially important for user interfaces
- ❑ *Resource Sharing* – threads share resources of process, easier than shared memory or message passing (IPC)
- ❑ *Economy* – cheaper than process creation, thread switching lower overhead than context switching
- ❑ *Scalability* – process can take advantage of multicore architectures




Multicore Programming

- ❑ *Multicore* or *multiprocessor systems* putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- ❑ *Parallelism* implies a system can perform more than one task simultaneously
- ❑ *Concurrency* supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

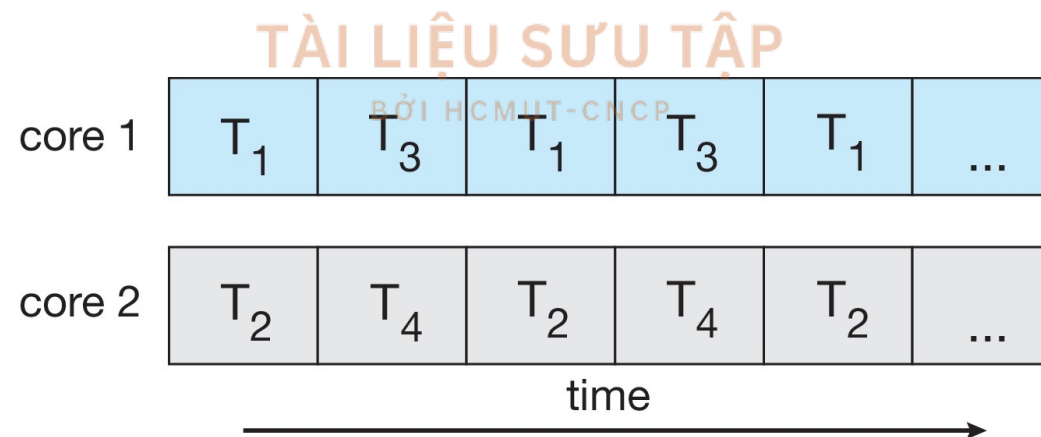


Concurrency vs. Parallelism

❑ Concurrent execution on single-core system:



❑ Parallelism on a multi-core system:



Multicore Programming


□ Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- **Task parallelism** – distributing threads across cores, each thread performing unique operation

TÀI LIỆU SƯU TẬP
BỞI HCMUT-CNCP



Data and Task Parallelism



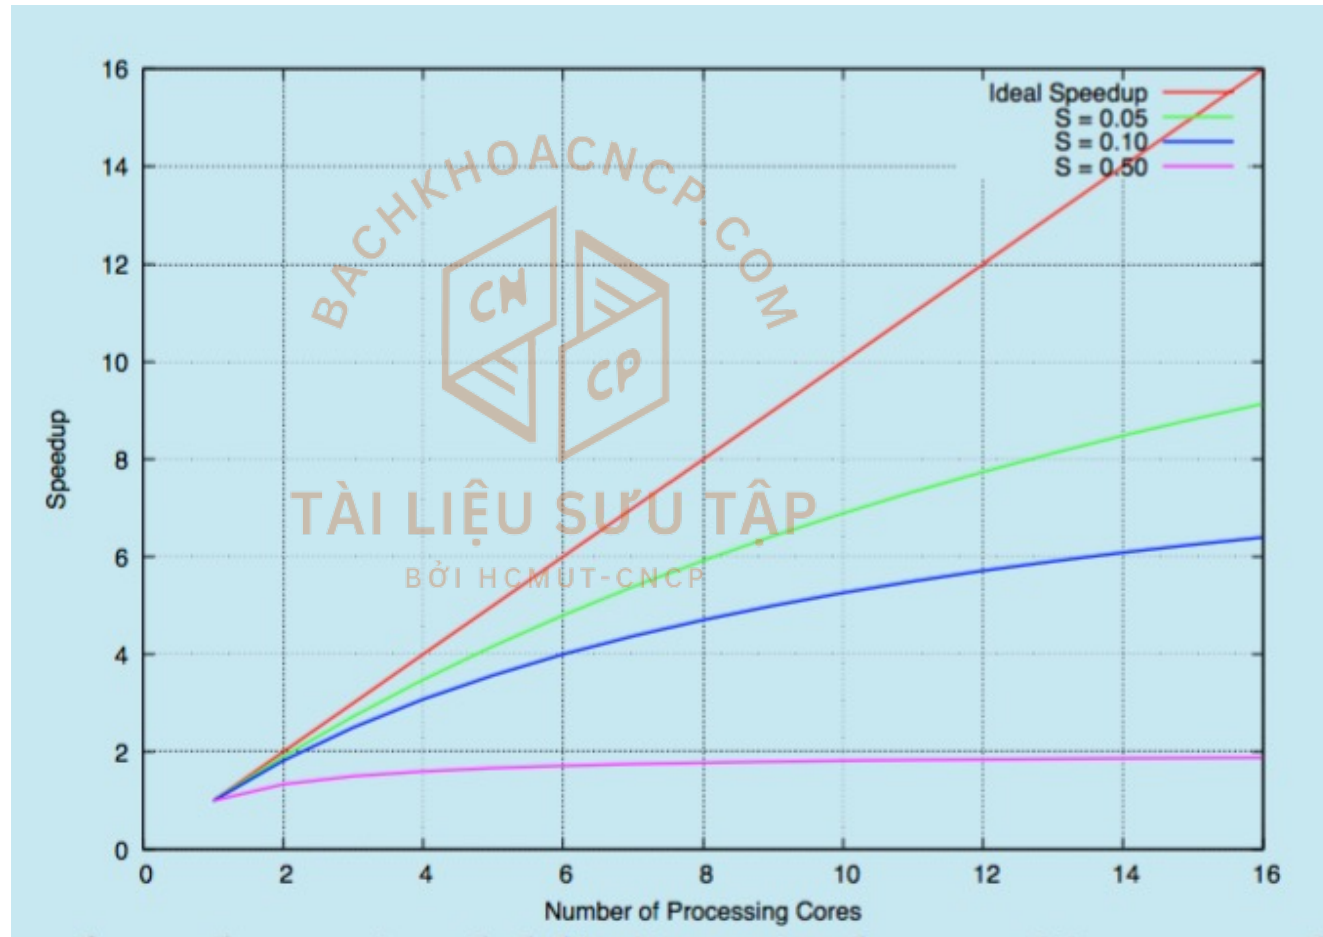
Amdahl's Law

- ❑ Identifies *performance gains* from adding additional cores to an application that has both serial and parallel components
 - S is serial portion
 - N processing cores
 - That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
 - As N approaches infinity, speedup approaches $1/S$
- ❑ Serial portion of an application has disproportionate effect on performance gained by adding additional cores
- ❑ But does the law take into account *contemporary multicore systems*?

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$



Amdahl's Law




User Threads and Kernel Threads

- ❑ *User threads* - management done by user-level threads library
- ❑ Three primary thread libraries:
 - POSIX Pthreads
 - Windows threads
 - Java threads
- ❑ *Kernel threads* - supported by the Kernel
- ❑ Examples – virtually all general purpose operating systems, including:
 - Windows, Linux, Mac OS X
 - iOS, Android



User and Kernel Threads




Multithreading Models

- ☐ Many-to-One
- ☐ One-to-One
- ☐ Many-to-Many



Many-to-One


- ❑ *Many user-level threads* mapped to *single kernel thread*
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- ❑ *Few systems currently use this model*
- ❑ Examples:
 - Solaris Green Threads
 - GNU Portable Threads



One-to-One


- ❑ *Each user-level thread* maps to *one kernel thread*
- ❑ Creating a user-level thread creates a kernel thread
- ❑ More concurrency than many-to-one
- ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Examples

- Windows
- Linux



Many-to-Many Model

- ❑ Allows *many user level threads* to be mapped to *many kernel threads*
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Windows with the **ThreadFiber** package
- ❑ Otherwise *not very common*



Thread Libraries

- ❑ **Thread library** provides programmer with API for creating and managing threads
- ❑ Two primary ways of implementing
 - Library entirely *in user space*
 - Kernel-level library *supported by the OS*

TÀI LIỆU SƯU TẬP
BỞI HCMUT-CNCP



Pthreads

- ❑ May be provided either as *user-level* or *kernel-level*
- ❑ A **POSIX standard (IEEE 1003.1c) API** for thread creation and synchronization
- ❑ Specification, not implementation
- ❑ API specifies behavior of the thread library, implementation is up to development of the library
- ❑ Common in UNIX operating systems (Linux & Mac OS X)



Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```



Pthreads Example (cont)

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```



Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10


/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



Implicit Threading

- ❑ Growing in popularity as numbers of threads increase, program correctness more difficult with *explicit threads*
- ❑ *Creation and management of threads done by compilers and run-time libraries* rather than programmers
- ❑ Five methods explored
 - Thread Pools
 - Fork-Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Threading Building Blocks




Thread Pools

- ❑ Create a *number of threads in a pool* where they await work
- ❑ Advantages:
 - Usually slightly *faster* to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be *bound* to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - ▶ i.e., Tasks could be scheduled to run periodically



Fork-Join Parallelism

- ❑ *Multiple threads (tasks) are forked, and then joined.*



Fork-Join Parallelism

- General algorithm for *fork-join strategy*:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))

    result1 = join(subtask1)
    result2 = join(subtask2)

  return combined results
```



Threading Issues

- ❑ Semantics of **fork()** and **exec()** system calls
- ❑ Signal handling
 - Synchronous and asynchronous
- ❑ Thread cancellation of target thread
 - Asynchronous or deferred
- ❑ Thread-local storage
- ❑ Scheduler Activations



Semantics of `fork()` and `exec()`

- ❑ Does `fork()` duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of `fork`
- ❑ `exec()` usually works as normal – replace the running process including all threads



Signal Handling

- ❑ **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- ❑ A **signal handler** is used to process signals
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Signal is handled by one of two signal handlers
 - ▶ default
 - ▶ user-defined
- ❑ Every signal has a default handler that kernel runs when handling signal
 - User-defined signal handler can override default
 - For single-threaded, signal delivered to process



Signal Handling (Cont.)

- ❑ *Where* should a signal be delivered for *multi-threaded*?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

TÀI LIỆU SƯU TẬP
BỞI HCMUT-CNCP



Thread Cancellation

- ❑ Terminating a thread before it has finished
- ❑ Thread to be canceled is target thread

- ❑ Two general approaches:

- Asynchronous cancellation terminates the target thread immediately
- Deferred cancellation allows the target thread to periodically check if it should be cancelled

```
pthread_t tid;
```

```
/* create the thread */
```

```
pthread_create(&tid, 0, worker, NULL);
```

```
/* cancel the thread */
```

```
pthread_cancel(tid);
```

```
/* wait for the thread to terminate */
```

```
pthread_join(tid, NULL);
```

- ❑ **Pthread** code to create and cancel a thread:



Thread Cancellation (Cont.)

- ❑ Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state


Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- ❑ If thread has cancellation disabled, cancellation remains pending until thread enables it
- ❑ Default type is deferred
 - Cancellation only occurs when thread reaches cancellation point
 - ▶ i.e., `pthread_testcancel()`
 - ▶ Then cleanup handler is invoked
- ❑ On Linux systems, thread cancellation is handled through signals



Scheduler Activations

- ❑ Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ❑ Typically use an intermediate data structure between user and kernel threads – *lightweight process* (LWP)
 - Appears to be a *virtual processor* on which process can schedule user thread to run
 - Each LWP attached to kernel thread
 - How many LWPs to create?
- ❑ Scheduler activations provide upcalls - a communication mechanism from the kernel to the upcall handler in the thread library
- ❑ This communication allows an application to maintain the correct number kernel threads



End of Chapter 4

