Chapter 8: Main Memory

- Background

- Contiguous Memory Allocation

- Paging

- Structure of the Page Table

- Swapping

- Example: The Intel 32 and 64-bit Architectures

- Example: ARMv8 Architecture

Objectives

- To provide a detailed description of various ways of *organizing memory hardware*

- To discuss various *memory-management techniques*

- To provide a detailed description of the **Intel Pentium**, which supports both *pure segmentation* and *segmentation with paging*

Background

- Program must be brought (from *disk*) into *memory* and placed within a process for it to be run

- *Main memory* and *registers* are only storage CPU can access directly
  - Register access is done in one CPU clock (or less)
  - Main memory access can take many cycles, causing a *stall*

- *Cache* sits between main memory and CPU registers

- *Memory Management Unit* (MMU) only sees a stream of:
  - addresses and *read* requests,
  - Or, address + data and *write* requests

- *Protection of memory* required to ensure correct operation

BACHKHOACNCP.COM

Protection

- Need to ensure that *a process can only access those addresses* in it address space.

- We can provide this protection by using a pair of *base* and *limit registers* define the contiguous *physical address space of a process*
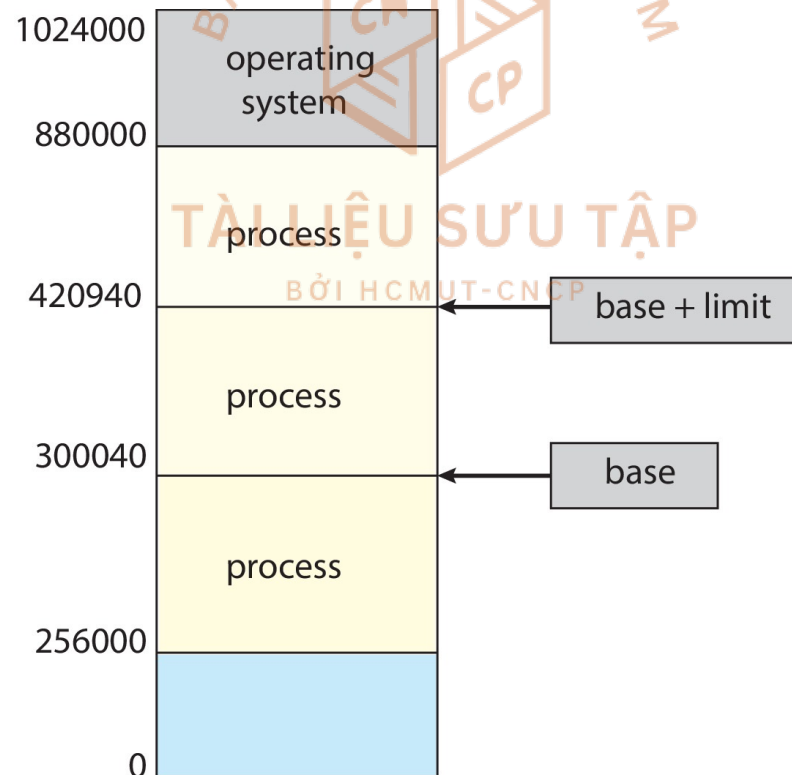
Hardware Address Protection

- CPU must *check every memory access* generated in user mode to be sure it is between (base) and (base + limit) for that user

- Instructions to loading the base and limit registers are *privileged*

Address Binding

- **Programs on disk**, ready to be brought into memory to execute, form an *input queue*

    - Without support, it must be loaded into address *0000*

    - Inconvenient to have first user process physical address always at *0000*

    - How can it not be?

- **Addresses** represented in different ways at different stages of a program's life

    - *Source code addresses* usually *symbolic*

    - *Compiled code addresses* bind to *relocatable addresses*

        ‣ e.g., "14 bytes from beginning of this module"

    - *Linker* or *loader* will bind relocatable addresses to *absolute addresses*

        ‣ e.g., 74014

   - Each binding maps one address space to another

Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages

  - *Compile time*: If memory location known a priori (e.g., Arduino), *absolute code* can be generated; must recompile code if starting location changes

  - *Load time*: Must generate *relocatable code* if memory location is not known at compile time

  - *Execution time*: Binding delayed until run time if the process can be moved from one memory segment to another during its execution

    - Need hardware support for address maps (e.g., base and limit registers)

Silberschatz, Galvin and Gagne ©2018

# Multistep Processing of a User Program
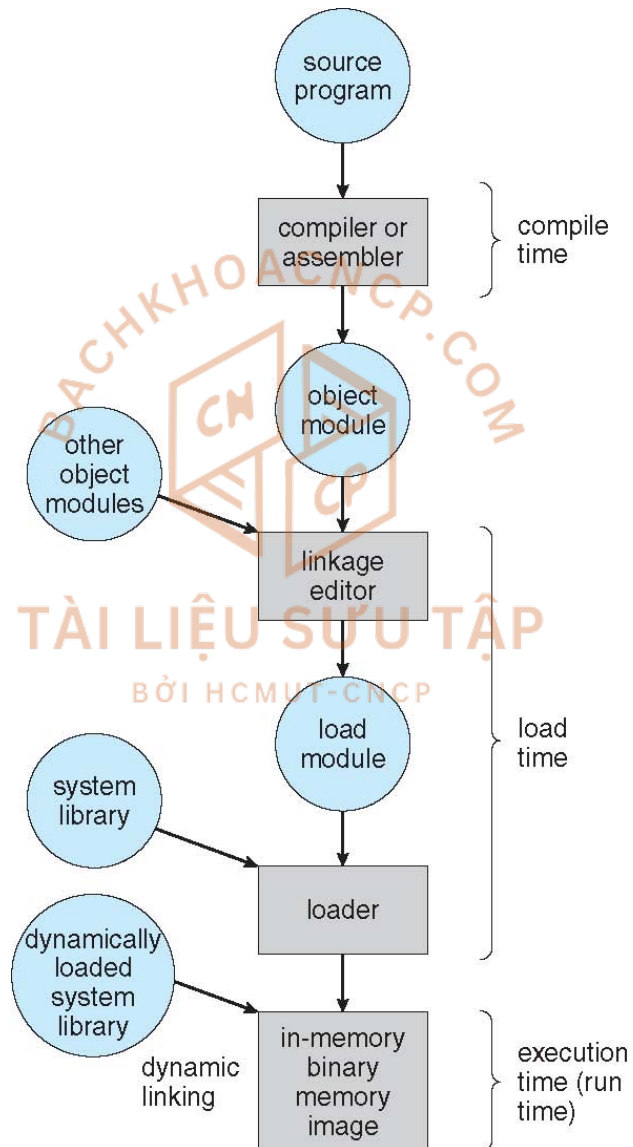
Logical vs. Physical Address Space

- The concept of a *logical address space* that is bound to a separate *physical address space* is central to proper memory management

    - *Logical address* – generated by the CPU; also referred to as *virtual address*

    - *Physical address* – address seen by the Memory Management Unit

- Logical and physical addresses are the same in *compile-time and load-time address-binding schemes*; logical (virtual) and physical addresses differ in *execution-time address-binding scheme*

    - *Logical address space* is the set of all logical addresses generated by a program

    - *Physical address space* is the set of all physical addresses generated by a program

Memory-Management Unit (MMU)

- *Hardware device* that maps virtual to physical address at run time
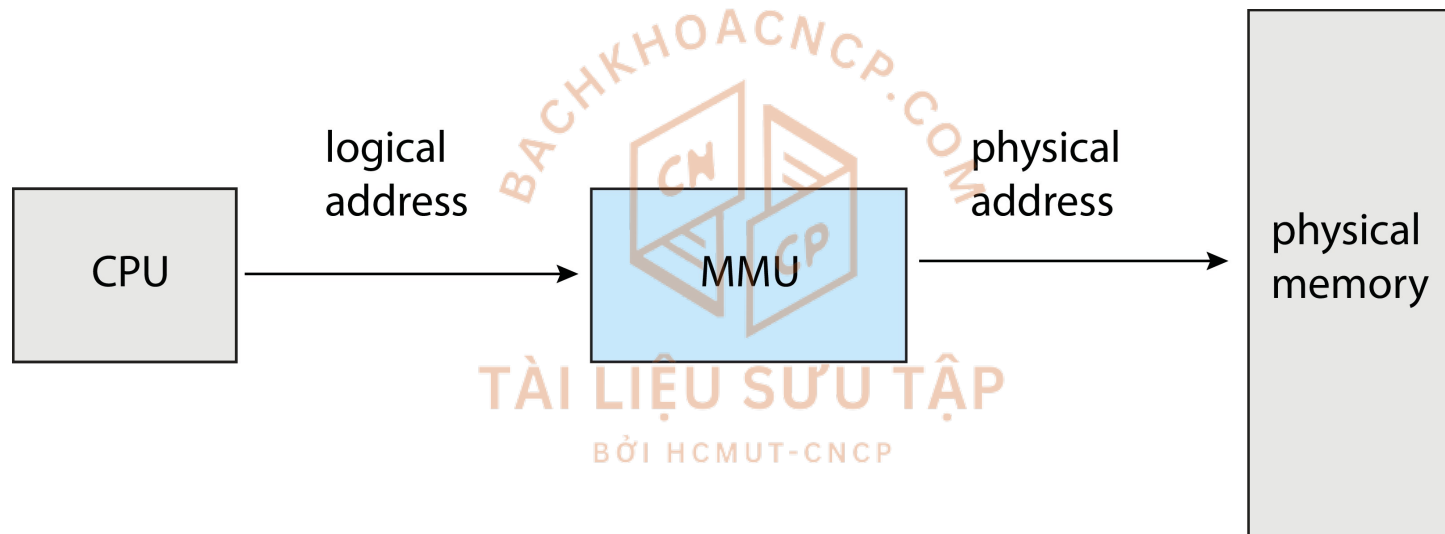
Memory-Management Unit (Cont.)

- E.g., Consider a simple scheme which is a generalization of the *base-register scheme*

  - The base register now called *relocation register*

  - The value in the relocation register is added to every address generated by a user process at the time it is sent to memory

- *The user program deals with logical addresses*; it never sees the real physical addresses

- *Execution-time binding* (i.e., logical addresses bound to physical addresses) occurs when reference is made to location in memory

Dynamic Loading

- The *entire program* does need to be in memory to execute

- *Routine* is not loaded until it is called

- Better *memory-space utilization*; unused routine is never loaded

- All routines kept on disk in *relocatable load format*

  - Useful when large amounts of code are needed to handle infrequently occurring cases

- No special support from the operating system is required

  - Implemented through program design

  - *OS can help by providing libraries to implement dynamic loading*

Dynamic Loading (Cont.)

- *Static loading* – system libraries and program code combined by the loader into the binary program image

- *Dynamic loading* – load postponed until execution time

- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine

  - Stub replaces itself with the address of the routine, and executes the routine

- Operating system checks if routine is in the memory address space of the process

  - If not in address space, add to address space

- *Dynamic loading is particularly useful for libraries*

  - System also known as *shared libraries*

  - Consider applicability to patching system libraries

    - Versioning may be needed

Dynamic linking (2)

- **static** linking

0x08048000

program

```
main:
...
call printf

printf:
...
ret
```

copy from libc

Dynamic linking (3)

- Dynamic linking

0x08048000

program

```
main:
  ...
  call printf
```

PLT
(r/o code)

```
printf:
  call GOT[5]
```

GOT
(r/w data)

```
...
[5]: dlfixup
...
```

```
dlfixup:
  GOT[5] = &printf
  call printf
```

0x40001234

libc

```
printf:
  ...
  ret
```

Contiguous Allocation

- *Main memory* must support both OS and user processes
  - Limited resource, must allocate efficiently

- *Contiguous allocation* is one early method
  - Main memory usually divides into *two partitions*:
    - ▸ *Resident operating system*, usually held in *high memory* with interrupt vector
    - ▸ *User processes* then held in *low memory*
      - – Each process contained in *single contiguous section of memory*

- *Relocation registers* (base registers) used to protect user processes from each other, and from changing OS code and data
  - *Base register* contains value of smallest physical address
  - *Limit register* contains range of logical addresses started from 0

- **MMU** maps logical address *dynamically*
  - Can then allow actions such as kernel code being *transient* and kernel changing *size*
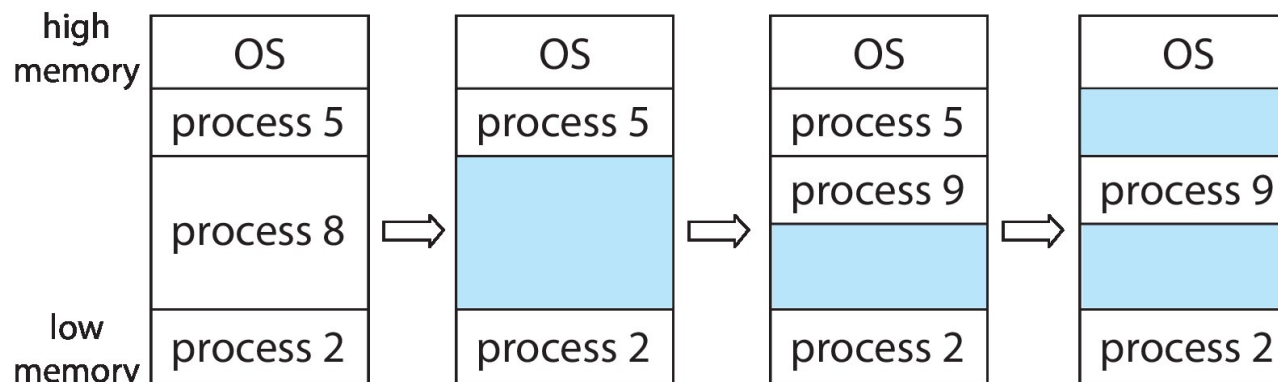
Hardware Support for Relocation and Limit Registers
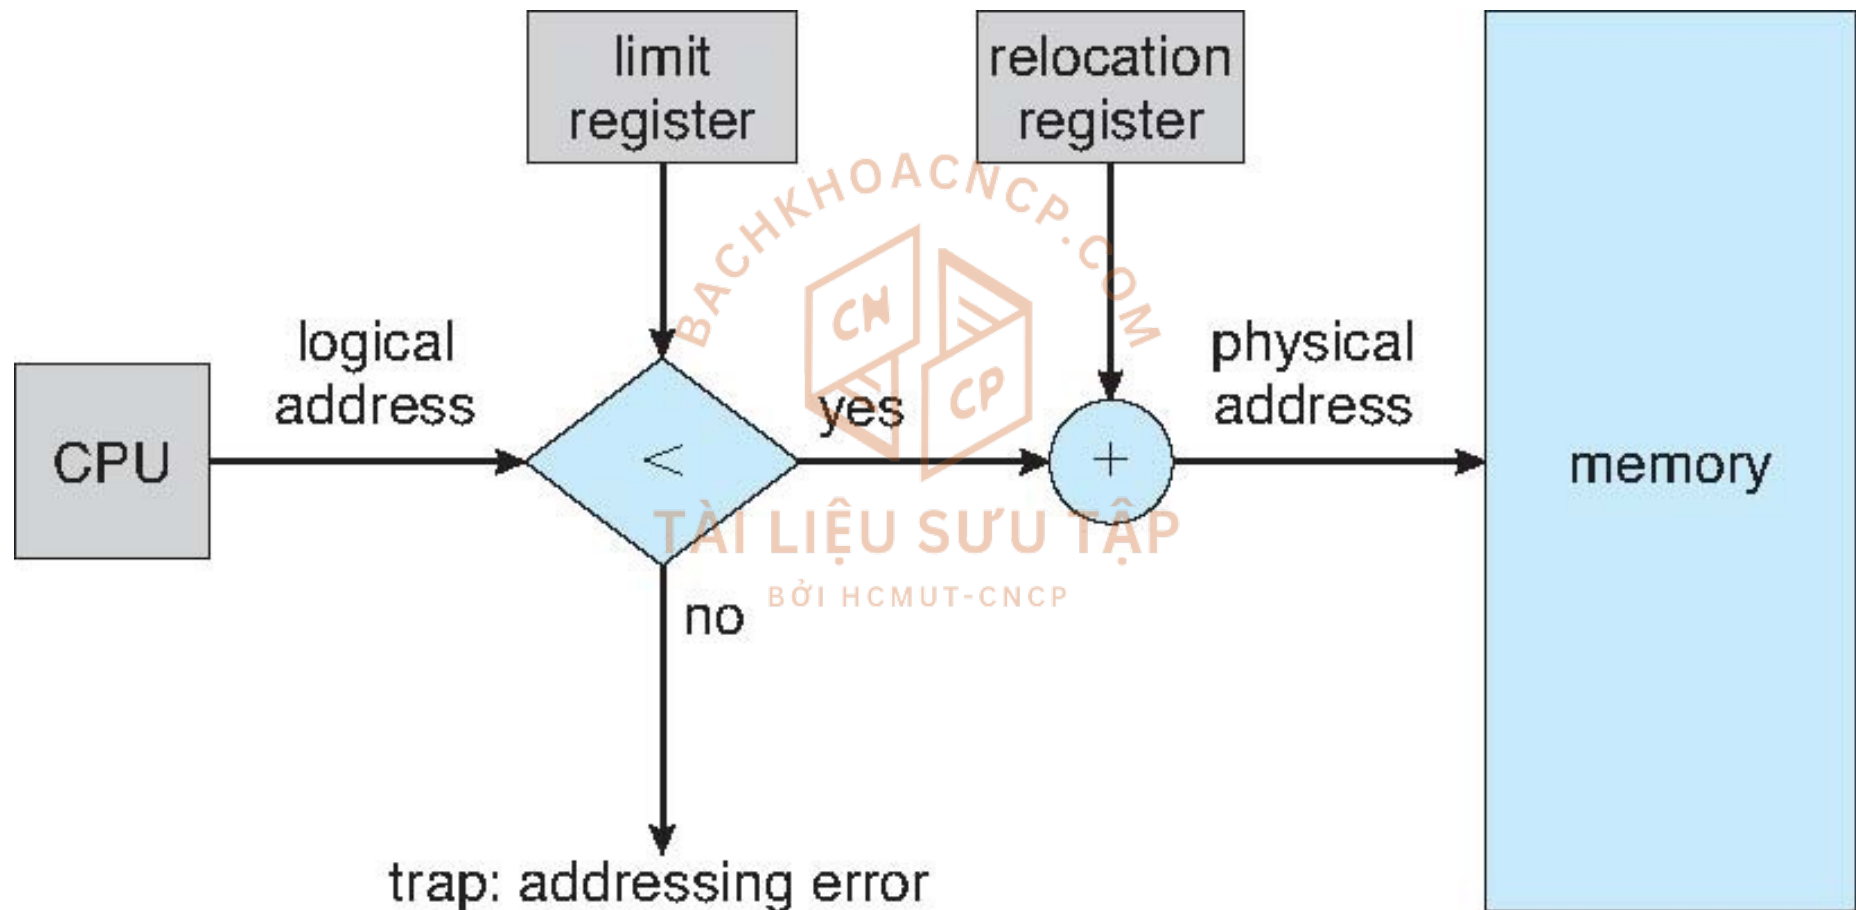
Variable Partition

- *Multiple-partition* allocation
  - *Degree of multiprogramming* limited by number of partitions
  - *Variable partition sizes* for efficiency (sized to the needs of a given process)
  - *Hole* – block of available memory;
    - holes of various size are scattered throughout memory
    - When a process arrives, it is allocated memory from a hole large enough to accommodate
  - Process exiting frees its partition, adjacent free partitions are combined
  - Operating system maintains information about: a) *allocated partitions* and b) *free partitions* (hole)

high memory	OS		OS		OS		OS
	process 5		process 5		process 5		
					process 9		process 9
	process 8						
low memory	process 2		process 2		process 2		process 2

BACHKHOACNCP.COM

Dynamic Memory-Allocation Problem

- How to satisfy a memory request of size *n* from a list of free holes?

  - *First-fit*: Allocate the *first* hole that is big enough

  - *Best-fit*: Allocate the *smallest* hole that is big enough

    - It must search entire list, unless ordered by size

    - Produces the smallest leftover hole

  - *Worst-fit*: Allocate the *largest* hole

    - It must also search entire list

    - Produces the largest leftover hole

- First-fit and best-fit strategies are better than worst-fit in terms of *speed* and *memory utilization*

■ Given five memory partitions of 100 KB, 500 KB, 200 KB, 300 KB, and 600 KB (in order), how would each of the first-fit, best-fit, worst-fit algorithms place processes of 212 KB, 417 KB, 112 KB, and 426 KB (in order)? Which algorithm makes the most efficient use of memory?

Fragmentation

- *External fragmentation* – total memory space exists to satisfy a request, but it is not contiguous

- *Internal fragmentation* – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used

- **Observation**

  - *First-fit* analysis reveals that given *N* blocks allocated, *0.5 N* blocks lost to fragmentation

  - *1/3* of that may be unusable => *50-percent rule*

Fragmentation (Cont.)

- Reduce external fragmentation by *compaction*

  - Shuffle memory contents to place all free memory together in one large block

  - Compaction is possible *only* if relocation is dynamic, and is done at execution time

  - I/O problem

    - Latch job in memory while it is involved in I/O

    - Do I/O only into OS buffers

- Now consider that **backing store** has same fragmentation problems

Paging

- *Physical address space of a process can be noncontiguous*; process is allocated physical memory whenever the latter is available
  - Avoids *external fragmentation*
  - Avoids problem of *varying sized memory chunks*
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is *power of 2*, between *512* bytes and *16* Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
  - To run a process of size *N* pages, need to find *N* free frames and load process
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have *internal fragmentation*

Address Translation Scheme

- Address generated by CPU is divided into:

  - *Page number* (*p*) – used as an index into a page table which contains base address of each page in physical memory

  - *Page offset* (*d*) – combined with base address to define the physical memory address that is sent to the Memory Management Unit

- For given logical address space 2^m and page size 2^n

  - Number of pages: $p = 2^m/2^n$
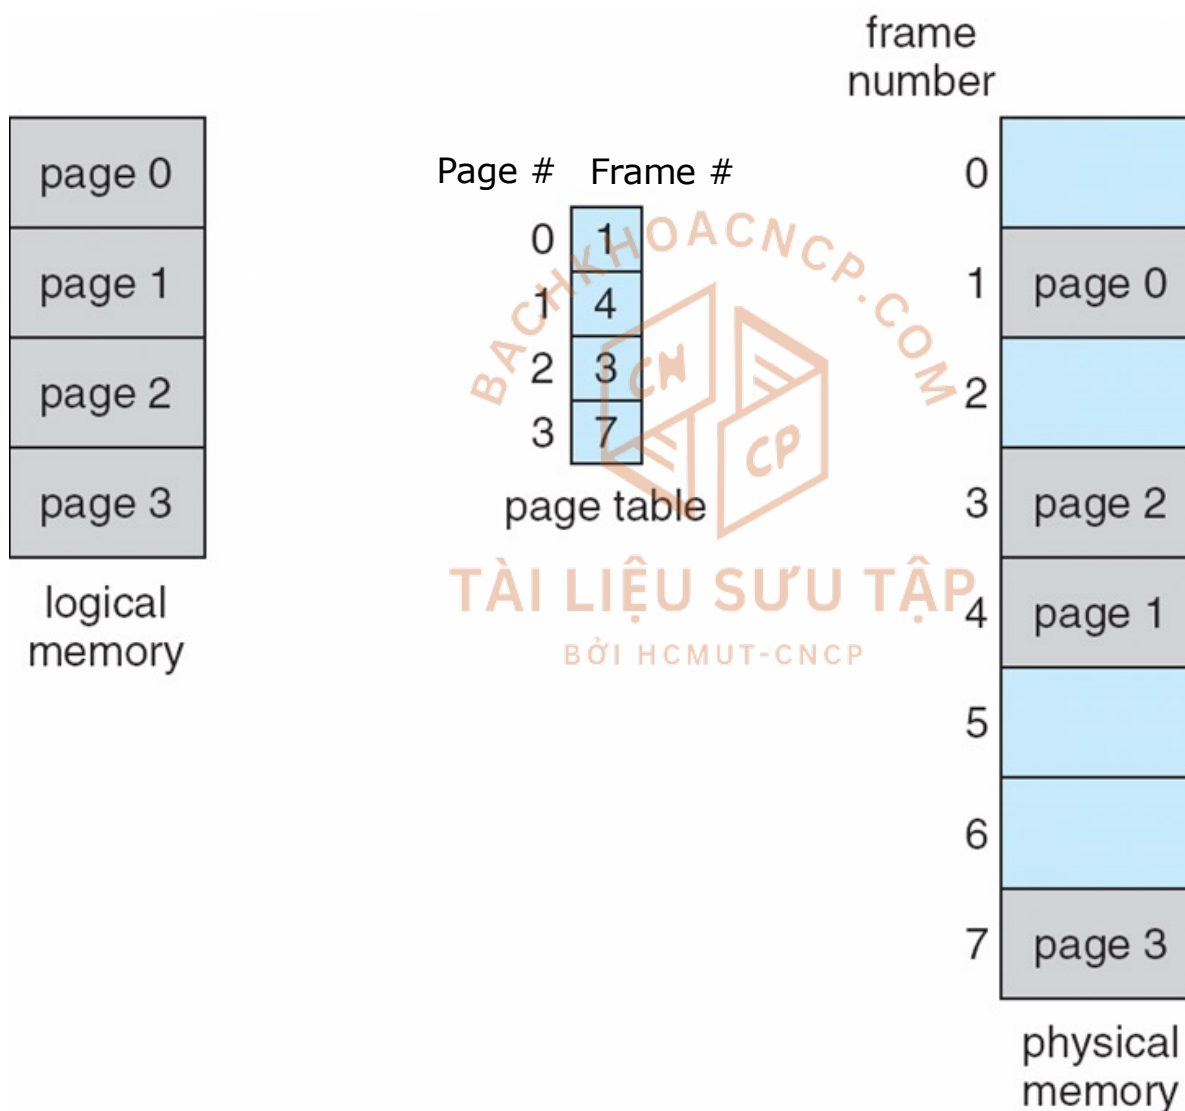
  - Offset inside a page: $d \in [0,n]$

page number	page offset
p	d
m -n	n
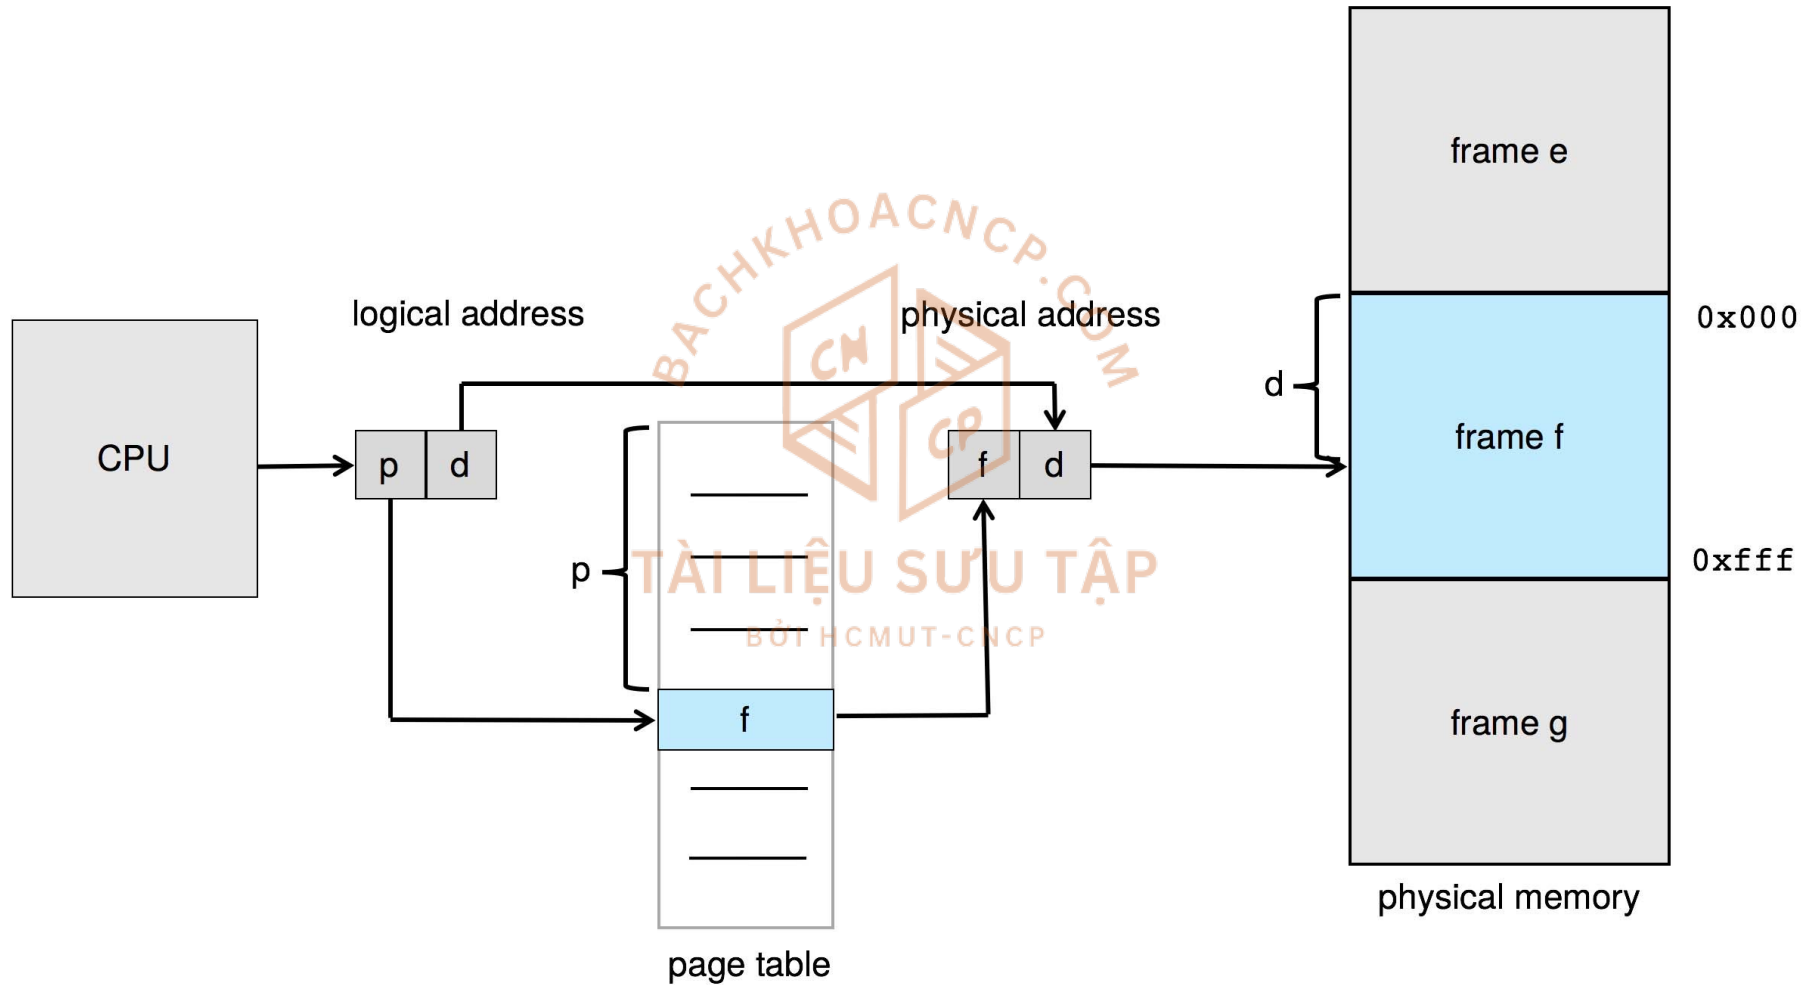
Paging Hardware

Example of Paging

- Using a page size of *4* bytes, a logical memory space of *16* bytes (*4* pages) and a physical memory of *32* bytes (*8* frames)

- Page size: 2^n
  - *n = 2*

- Logical address space: 2^m
 - *m = 4*

- Physical address space: 2^r
  - *r = 5*

Paging – Calculating internal fragmentation

- Page size = *2,048* bytes

- Process size (space) = *72,766* bytes

    - *35* pages + *1,086* bytes

- Internal fragmentation of last page = *2,048 - 1,086* = *962* bytes

    - Worst case fragmentation = 1 frame – 1 byte

    - On average fragmentation = 1 / 2 frame size

- So small frame sizes desirable?

    - But each page table entry takes memory to track

- Page sizes growing over time

    - Solaris supports two page sizes – **8** KB and **4** MB

Free Frames

(a) Before allocation

(b) After allocation

Implementation of Page Table

- Page table is kept in *main memory*

  - *Page-table base register* (**PTBR**) points to the page table

  - *Page-table length register* (**PTLR**) indicates size of the page table

- In this scheme every data/instruction access requires *two memory accesses*

  - One for the *page table*

  - one for the *data / instruction*

- The two-memory-access problem can be solved by the use of a special fast-lookup hardware cache called *Translation Look-aside Buffers* (**TLBs**) (also called *associative memory*)

Translation Look-Aside Buffer (TLB)

- Some TLBs store *address-space identifiers* (**ASIDs**) in each TLB entry – uniquely identifies each process to provide address-space protection for that process

    - Otherwise need to *flush at every context switch*

- TLBs typically small (*64* to *1,024* entries)

- On a TLB miss, value is loaded into the TLB for faster access next time

    - *Replacement policies* must be considered

    - Some entries can be *wired down* for permanent fast access
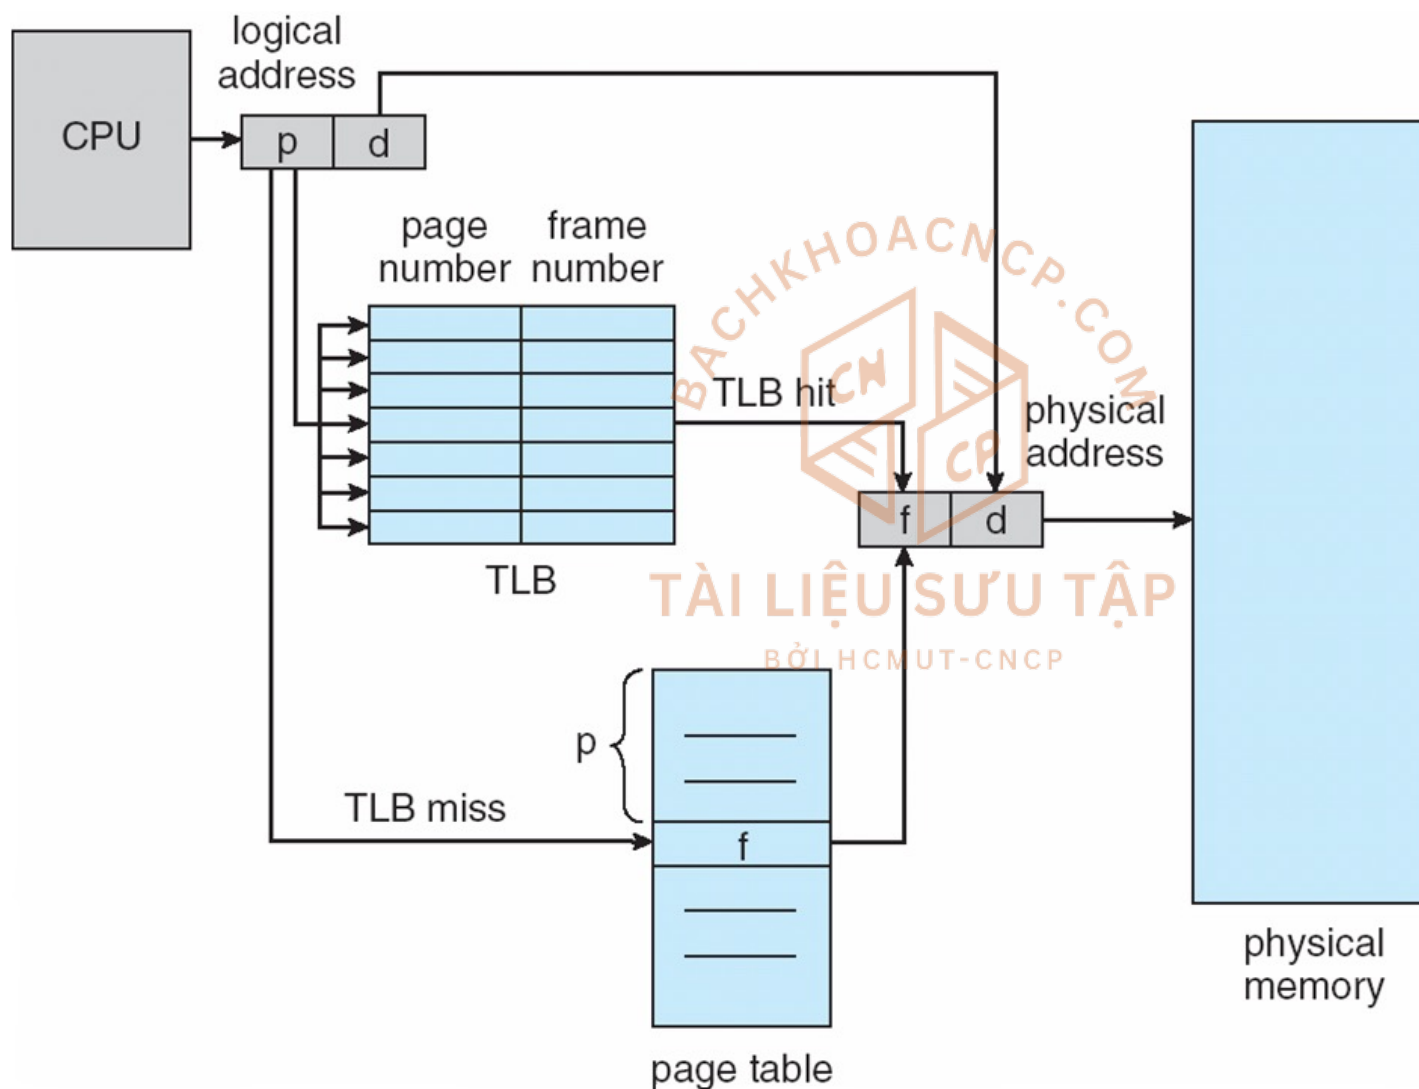
Hardware

- *Associative memory – parallel search*

Page #	Frame #

- *Address translation (p, d)*

  - If **p** is in associative register, get *frame #* out

  - Otherwise *get frame # from page table in memory*

Effective Access Time

- **Hit ratio** – percentage of times that a page number is found in TLB

    - E.g., An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.

- Suppose that *10 nanoseconds (ns)* to access memory.

    - If we find desired page in TLB then mapped-memory access take 10 ns

    - Otherwise we need *two memory access,* so it is 20 ns

- *Effective Access Time* (**EAT**)

    - EAT = 0.80 x 10 + 0.20 x 20 = *12 ns* => implying *20%* slowdown in access time

    - Consider a more realistic hit ratio of *99%*, EAT = 0.99 x 10 + 0.01 x 20 = *10.1 ns* => implying only *1%* slowdown in access time

Memory Protection

- Memory protection implemented by associating **protection bit** with each frame to indicate if *read-only* or *read-write* access is allowed

  - Can also add more bits to indicate page execute-only, and so on

- **Valid-invalid bit** attached to each entry in the page table:

  - "*valid*" indicates that the associated page is in the process' logical address space, and is thus a legal page

  - "*invalid*" indicates that the page is not in the process' logical address space

  - Or use *page-table length register* (**PTLR**)

- Any violations result in a *trap* to the kernel

Shared Pages

- *Shared code*

  - One copy of read-only (*reentrant*) code shared among processes (i.e., text editors, compilers, window systems)

  - Similar to *multiple threads* sharing the same process space

  - Also useful for *inter-process communication* if sharing of read-write pages is allowed

- *Private code and data*

  - Each process keeps a separate copy of the code and data

  - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example
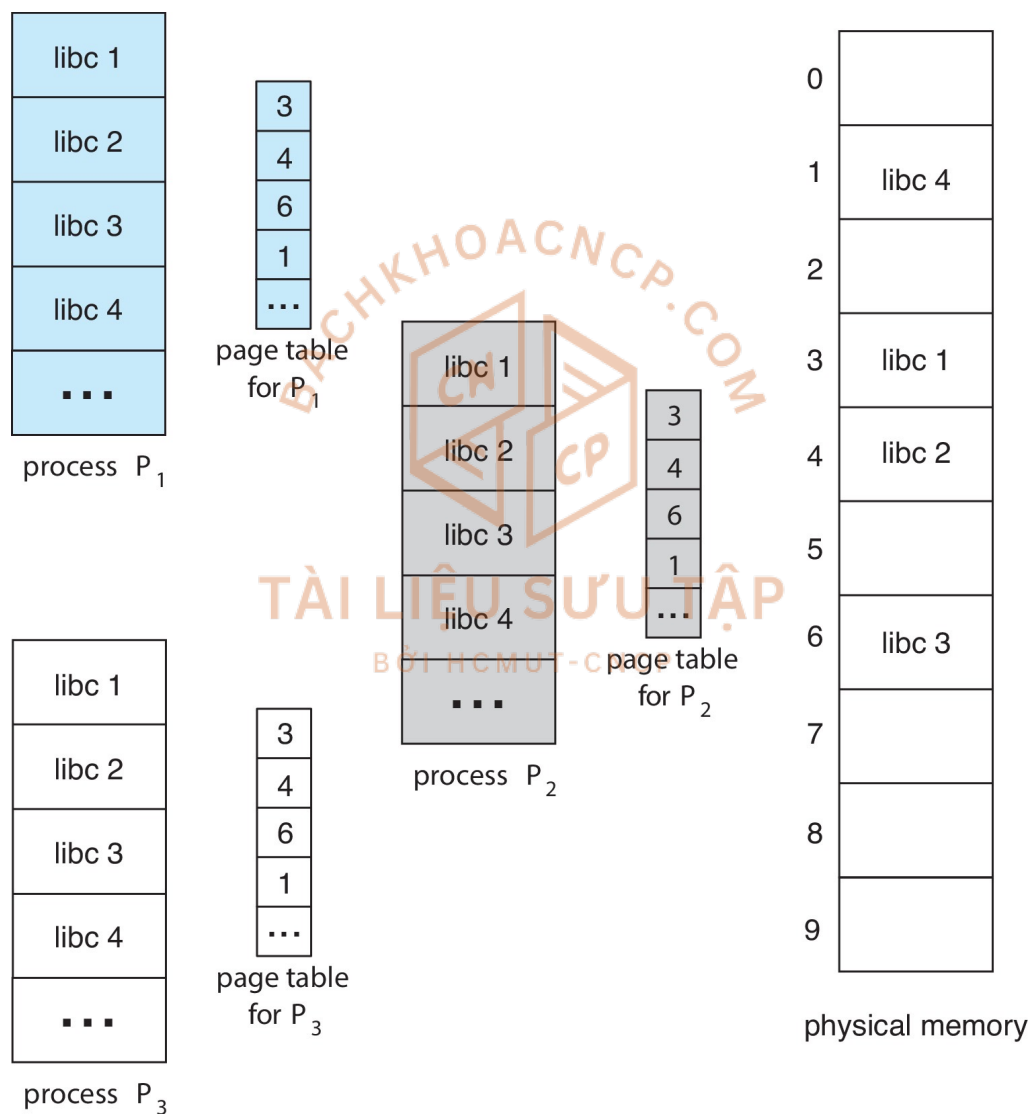
Structure of the Page Table

- *Memory structures* for paging can get huge using straight-forward methods

  - Consider a *32-bit logical address space* as on modern computers

  - Page size of **4** KB (2^{12})

  - Page table would have 1 million entries (2^{32} / 2^{12})

  - If each entry is 4 bytes => each process 4 MB of physical address space for the page table alone

    ▸ Don't want to allocate that contiguously in main memory

  - One simple solution is to divide the page table into smaller units

    ▸ *Hierarchical Paging*

    ▸ *Hashed Page Tables*

    ▸ *Inverted Page Tables*

Hierarchical Page Tables

- Break up the logical address space into multiple page tables

- A simple technique is a two-level page table

- We then page the page table

Silberschatz, Galvin and Gagne ©2018

Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:

  - a page number consisting of 22 bits

  - a page offset consisting of 10 bits

- Since the page table is paged, the page number is further divided into:

  - a 10-bit page number

  - a 12-bit page offset

page number		page offset
p_1	p_2	d
10	10	12

- Thus, a logical address is as follows:

  where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table

- Known as *forward-mapped page table*

Address-Translation Scheme

64-bit Logical Address Space

■ Even two-level paging scheme not sufficient

■ If page size is 4 KB (2^{12})

outer page	inner page	offset
p_1	p_2	d
42	10	12

- Then page table has 2^{52} entries

- If two level scheme, inner page tables could be 2^{10} 4-byte entries

- Address would look like

- Outer page table has 2^{42} entries or 2^{44} bytes

- One solution is to add a 2^{nd} outer page table

- But in the following example the 2^{nd} outer page table is still 2^{34} bytes in size

    ▸ And possibly 4 memory access to get to one physical memory location

outer page	inner page	offset
p_1	p_2	d
42	10	12

| 2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Hashed Page Tables

- Common in address spaces *> 32 bits*

- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location

- Each element contains (1) the *virtual page number* (2) the value of the *mapped page frame* (3) a *pointer* to the next element

- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted

- Variation for 64-bit addresses is *clustered page tables*
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for *sparse* address spaces (where memory references are non-contiguous and scattered)

Hashed Page Table
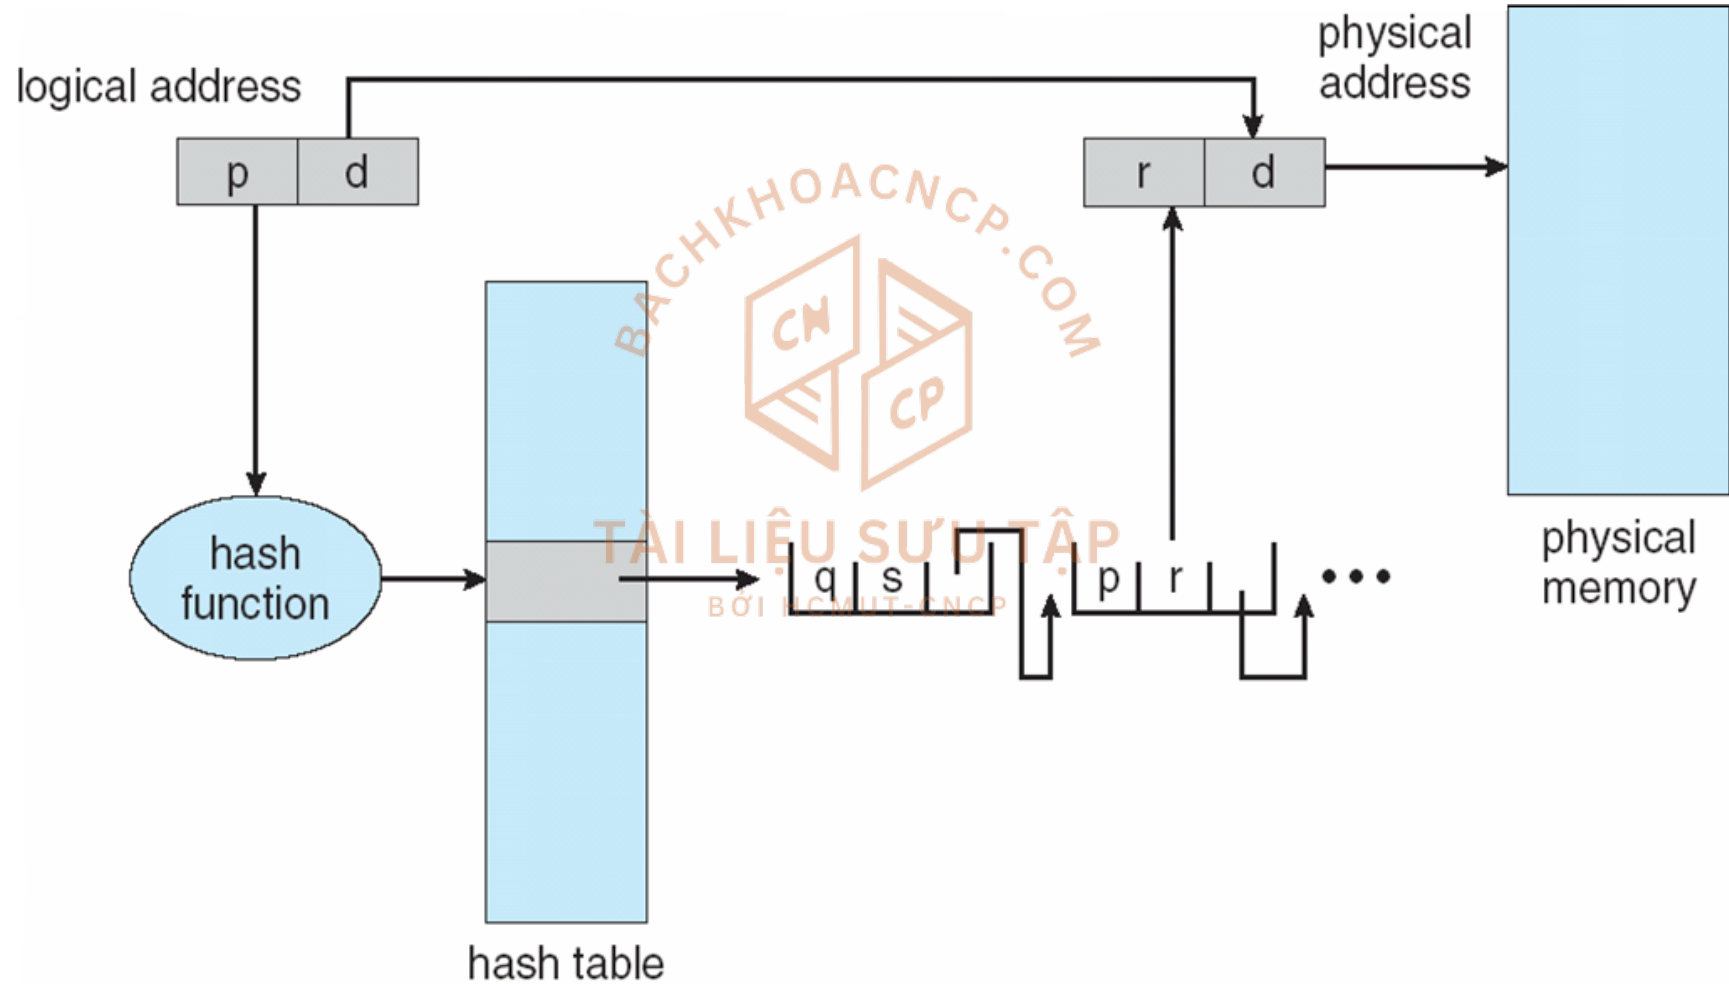
Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, *track all physical pages*

- One entry for each real page of memory

- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page

- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs

- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access

- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address

Inverted Page Table Architecture

■ Assume that a task is divided into four equal-sized segments and that the system builds an eight-entry page descriptor table for each segment. Thus, the system has a combination of segmentation and paging. Assume also that the page size is 2 Kbytes.

1. What is the maximum size of each segment?

2. What is the maximum logical address space for the task?

3. Assume that an element in physical location 00021ABC is accessed by this task. What is the format of the logical address that the task generates for it? What is the maximum physical address space for the system?

Exercise 3

- This question refers to an architecture using segmentation with paging. In this architecture, the 32-bit virtual address is divided into fields as follows:

    4 bit segment number

    12 bit page number

    16 bit offset

    Find the physical address corresponding to each of the following virtual addresses (answer "bad virtual address" if the virtual address is invalid).

    1. 00020000

    2. 20011002

    3. 10001111

Exercise 3 (cont.)

	Segment table			Page table A			Page table B
0	Page table A		0	CAFE		0	F000
1	Page table B		1	DEAD		1	D8BF
x	(rest invalid)		2	BEEF		X	(rest invalid)
			3	BA11			
			X				

Oracle SPARC Solaris

- Consider modern, 64-bit operating system example with tightly integrated HW

  - Goals are efficiency, low overhead

- Based on hashing, but more complex

- Two hash tables

  - One kernel and one for all user processes

  - Each maps memory addresses from virtual to physical memory

  - Each entry represents a contiguous area of mapped virtual memory,

    ▸ More efficient than having a separate hash-table entry for each page

  - Each entry has base address and span (indicating the number of pages the entry represents)

Oracle SPARC Solaris (Cont.)

- TLB holds translation table entries (TTEs) for fast hardware lookups

  - A cache of TTEs reside in a translation storage buffer (TSB)

    - Includes an entry per recently accessed page

- Virtual address reference causes TLB search

  - If miss, hardware walks the in-memory TSB looking for the TTE corresponding to the address

    - If match found, the CPU copies the TSB entry into the TLB and translation completes

    - If no match found, kernel interrupted to search the hash table

      - The kernel then creates a TTE from the appropriate hash table and stores it in the TSB, Interrupt handler returns control to the MMU, which completes the address translation.

Swapping

- A process can be *swapped* temporarily out of memory to a backing store, and then brought back into memory for continued execution

  - Total physical memory space of processes can exceed physical memory

- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images

- *Roll out, roll in* – swapping variant used for *priority-based scheduling* algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed

- Major part of swap time is *transfer time*; total transfer time is directly proportional to the amount of memory swapped

- System maintains a *ready queue* of ready-to-run processes which have memory images on disk
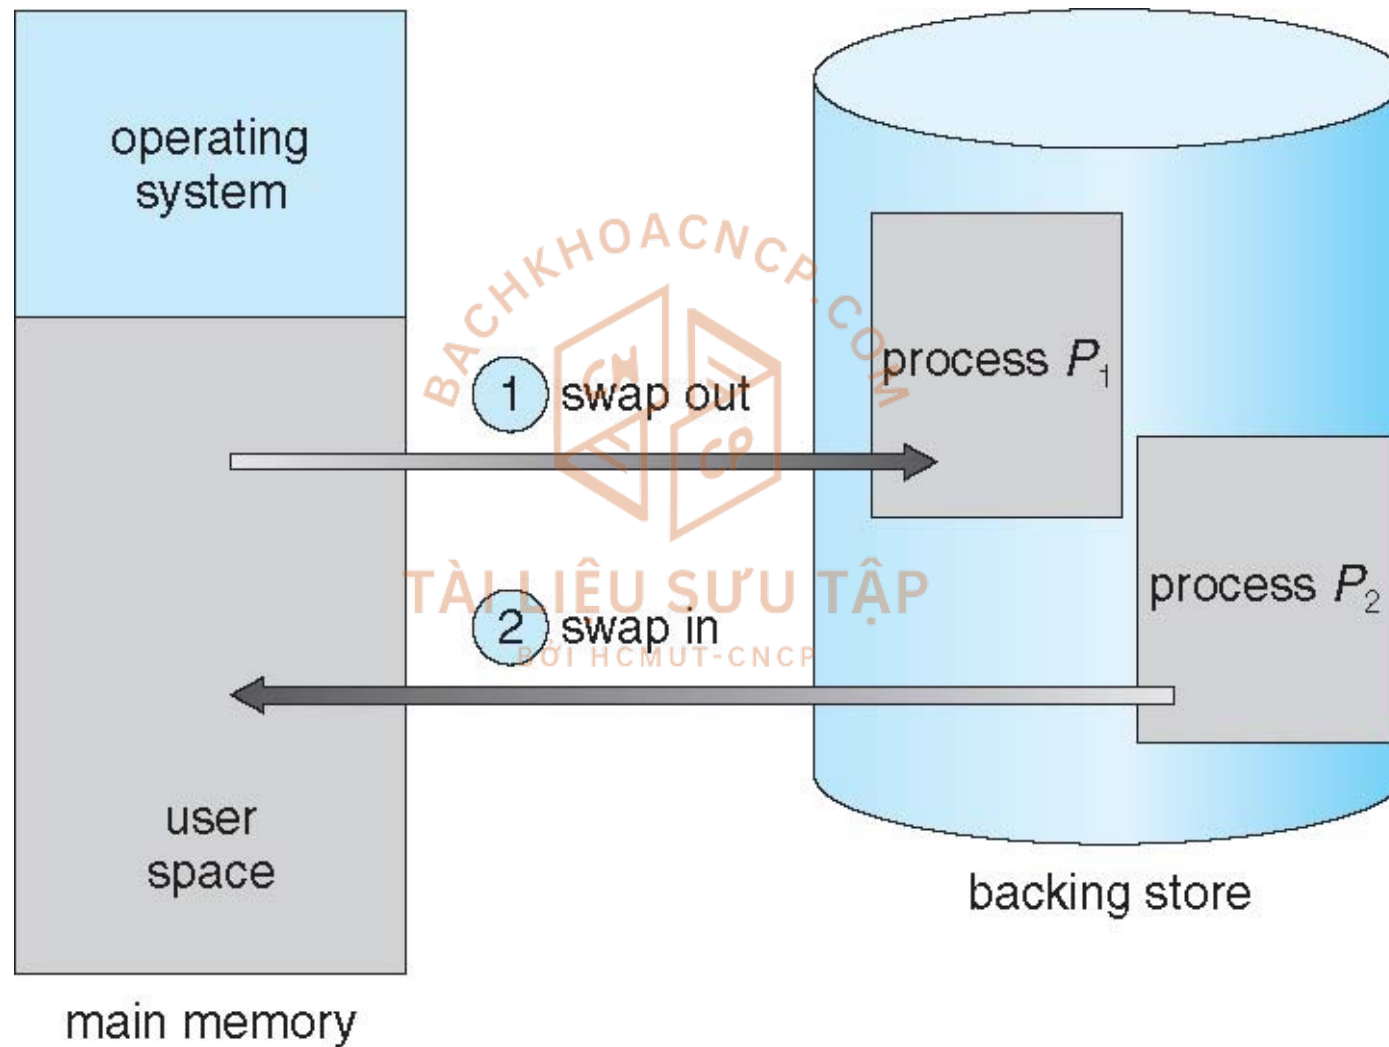
Swapping (Cont.)

- Does the swapped out process need to swap back in to same physical addresses?

- Depends on address binding method

  - Plus consider pending I/O to / from process memory space

- Modified versions of swapping are found on many systems (i.e., **UNIX**, **Linux**, and **Windows**)

  - Swapping normally disabled

  - Started if more than threshold amount of memory allocated

  - Disabled again once memory demand reduced below threshold

main memory

backing store

Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process

- Context switch time can then be very high

- 100MB process swapping to hard disk with transfer rate of 50MB/sec

  - Swap out time of 2000ms

  - Plus swap in of same sized process

  - Total context switch swapping component time of 4000ms (4 seconds)

- Can reduce if reduce size of memory swapped – by knowing how much memory really being used

  - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

- Other *constraints* as well on swapping

  - Pending I/O – can't swap out as I/O would occur to wrong process

  - Or always transfer I/O to kernel space, then to I/O device

    ▸ Known as *double buffering*, adds overhead

- Standard *swapping* not used in modern operating systems

  - But modified version common

    ▸ Swap only when free memory extremely low

Swapping on Mobile Systems

- **Not typically supported**

  - Flash memory based

    ▸ Small amount of space

    ▸ Limited number of write cycles

    ▸ Poor throughput between flash memory and CPU on mobile platform

- **Instead use other methods to free memory if low**

  - iOS asks apps to voluntarily relinquish allocated memory

    ▸ Read-only data thrown out and reloaded from flash if needed

    ▸ Failure to free can result in termination

  - Android terminates apps if low free memory, but first writes application state to flash for fast restart

  - Both OSes support paging as discussed below

Swapping with Paging

main memory

backing store

Example: The Intel 32 and 64-bit Architectures

- Dominant industry chips

- Pentium CPUs are 32-bit and called IA-32 architecture

- Current Intel CPUs are 64-bit and called IA-64 architecture

- Many variations in the chips, cover the main ideas here

Example: The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging

  - Each segment can be 4 GB

  - Up to 16 K segments per process

  - Divided into two partitions

    ▸ First partition of up to 8 K segments are private to process (kept in *local descriptor table* (**LDT**))

   ▸ Second partition of up to 8K segments shared among all processes (kept in *global descriptor table* (**GDT**))

■ **CPU generates logical address**

- Selector given to segmentation unit

  ▸ Which produces linear addresses

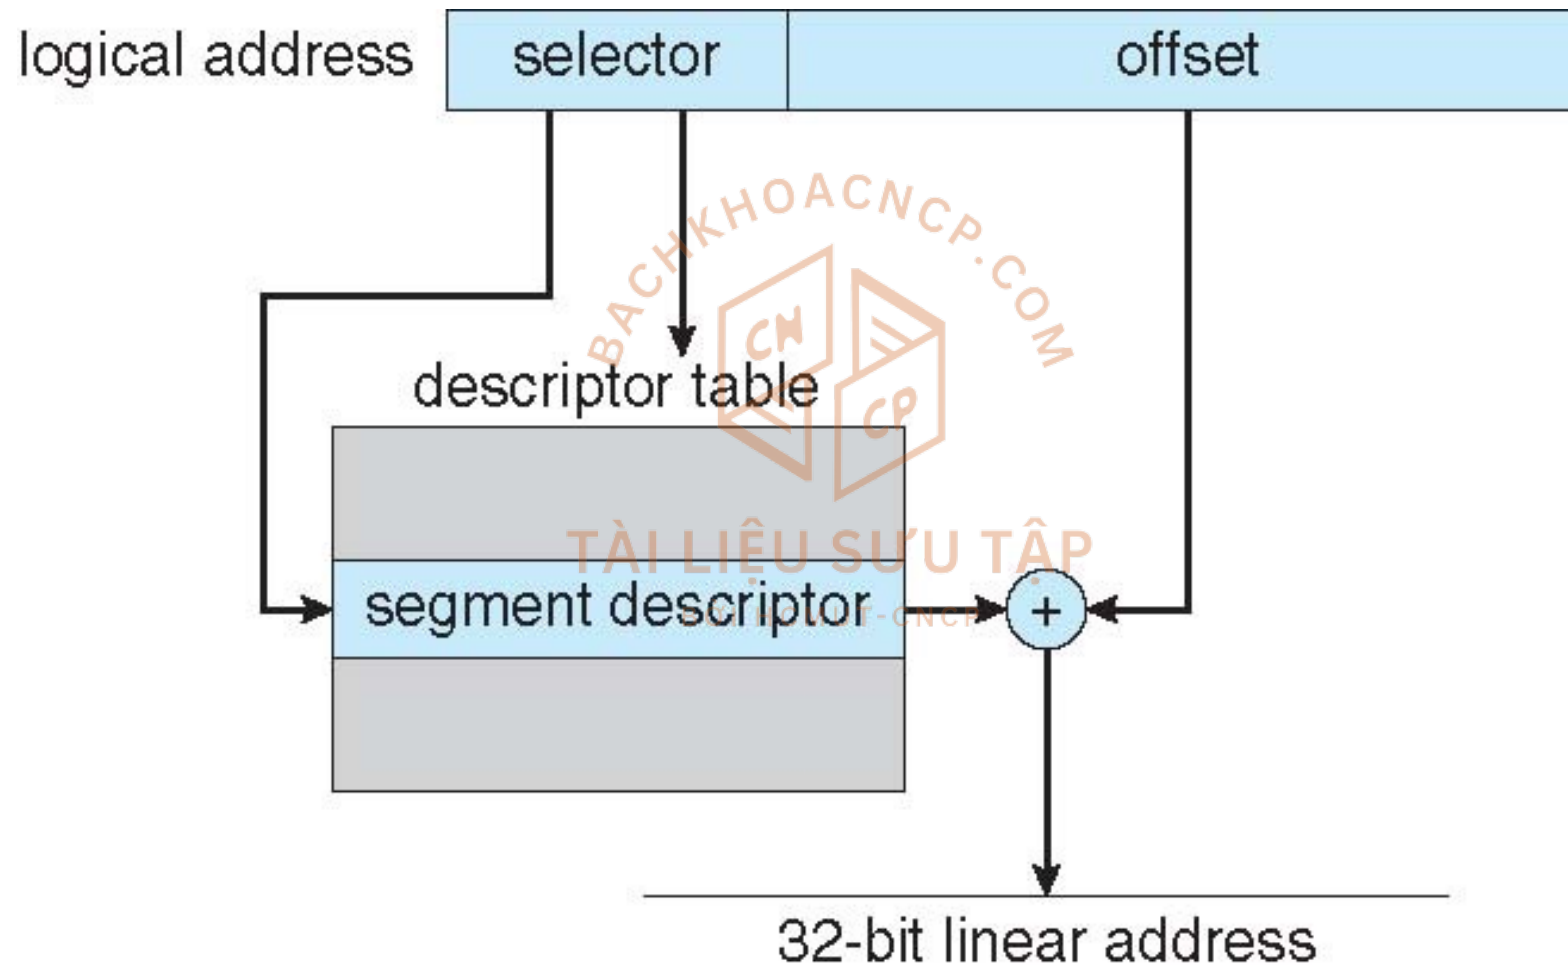| s	g	p
13	1	2

- Linear address given to paging unit

  ▸ Which generates physical address in main memory

  ▸ Paging units form equivalent of MMU

  ▸ Pages sizes can be 4 KB or 4 MB

| page number		page offset
p_1	p_2	d
10	10	12

Intel IA-32 Segmentation

Intel IA-32 Paging Architecture

Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create **page address extension** (**PAE**), allowing 32-bit apps access to more than 4GB of memory space

  - Paging went to a 3-level scheme

  - Top two bits refer to a **page directory pointer table**

  - Page-directory and page-table entries moved to 64-bits in size

  - Net effect is increasing address space to 36 bits – 64GB of physical
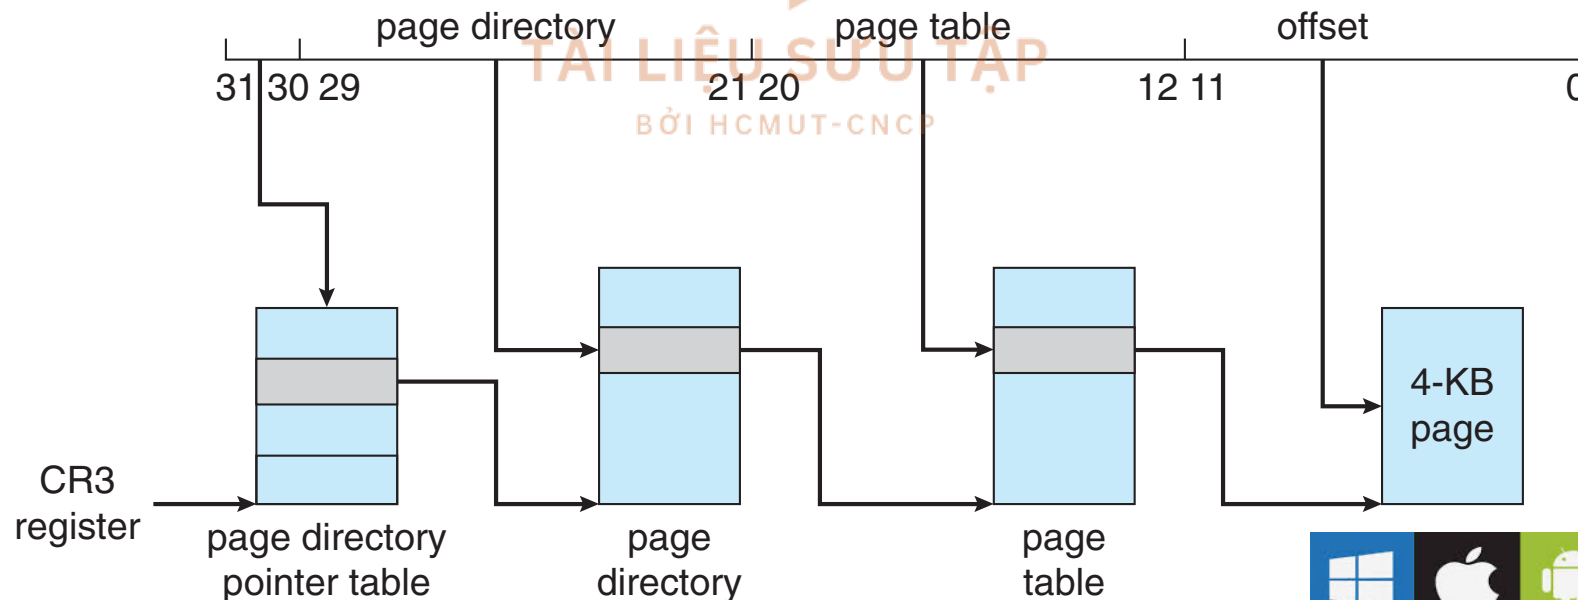
Intel x86-64

- Current generation Intel x86 architecture

- 64 bits is ginormous (> 16 exabytes)

- In practice only implement 48 bit addressing

  - Page sizes of 4 KB, 2 MB, 1 GB

  - Four levels of paging hierarchy

- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits

unused	page map level 4	page directory pointer table	page directory	page table	offset
63 48	47 39	38 30	29 21	20 12	11 0

Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)

- Modern, energy efficient, 32-bit CPU

- 4 KB and 16 KB pages

- 1 MB and 16 MB pages (termed *sections*)

- One-level paging for sections, two-level for smaller pages

- Two levels of TLBs

  - Outer level has two micro TLBs (one data, one instruction)

  - Inner is single main TLB

  - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU

Summary

- *Memory* is central to the operation of a modern computer system and consists of a large array of *bytes*, each with *its own address*.

- One way to *allocate an addresses space* to each process is through the use of base and limit registers. The *base register* holds the smallest legal physical memory address, and the *limit register* specifies the size of the range.

- *Binding* symbolic address references to actual physical addresses may occur during (1) *compile*, (2) *load*, or (3) *execution time*.

- An address generated by the CPU is known as a logical address, which the *Memory Management Unit* (**MMU**) *translates* to a physical address in memory.

Summary (Cont.)

- One approach to allocating memory is to *allocate partitions of contiguous memory of varying sizes*. These partitions may be allocated based on three possible strategies: (1) *first fit*, (2) *best fit*, and (3) *worst fit*.

- Modern operating systems use *paging* to manage memory. In this process, physical memory is divided into fixed-sized blocks called *frames* and logical memory into blocks of the same size called *pages*

- When paging is used, a logical address is divided into two parts: a page number and a page offset. The page number serves as an index into a *per-process page table* that contains the frame in physical memory that holds the page. The offset is the specific location in the frame being referenced.

Summary (Cont.)

■ A *Translation Look-aside Buffer* (**TLB**) is a hardware *cache* of the page table. Each TLB entry contains a page number and its corresponding frame.

■ Using a TLB in *address translation for paging systems* involves obtaining the page number from the logical address and checking if the frame for the page is in the TLB. If it is, the frame is obtained from the TLB. If the frame is not present in the TLB, it must be retrieved from the page table.

■ *Hierarchical paging* involves dividing a logical address into multiple parts, each referring to different levels of page tables. As addresses expand beyond 32 bits, the number of hierarchical levels may become large. Two strategies that address this problem are *hashed page tables* and *inverted page tables*.

Summary (Cont.)

- *Swapping* allows the system to move pages belonging to a process to disk to increase the *degree of multiprogramming*.

- The *Intel 32-bit architecture* has two levels of page tables and supports either 4-KB or 4-MB page sizes. This architecture also supports page- address extension, which allows 32-bit processors to access a physical address space larger than 4 GB. The *x86-64 and ARMv9 architectures* are 64-bit architectures that use hierarchical paging.

End of Chapter 9