


Artificial Intelligence

Q learning

TÀI LIỆU SƯU TẬP
BỞI HCMUT-CNCP


Pham Viet Cuong
Dept. Control Engineering & Automation, FEEE
Ho Chi Minh City University of Technology

- ✓ Supervised learning: Classification, regression
- ✓ Unsupervised learning: Clustering
- ✓ Reinforcement learning:
 - ❖ More general than supervised/unsupervised learning
 - ❖ Learn from interactive with environment (perform actions and observe rewards) to achieve a goal
 - ❖ Goal: Learn a policy to maximize some measure of long-term reward

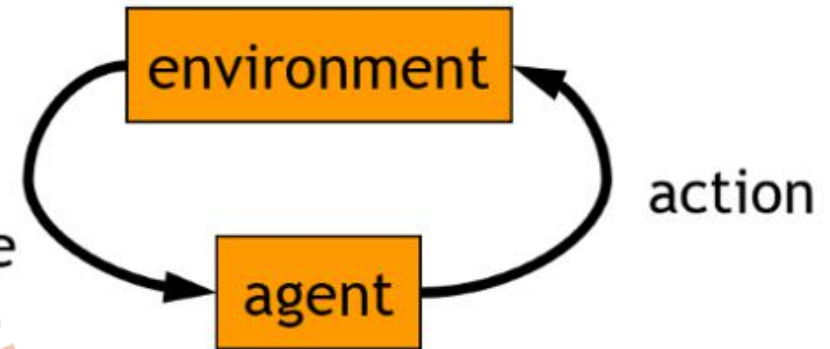


✓ Examples:

Cart-Pole Problem



reward
new state



Objective: Balance a pole on top of a movable cart



State: angle, angular speed, position, horizontal velocity

Action: horizontal force applied on the cart

Reward: 1 at each time step if the pole is upright

✓ Examples:

Robot Locomotion



Objective: Make the robot move forward

State: Angle and position of the joints

Action: Torques applied on joints

Reward: 1 at each time step upright + forward movement

✓ Examples: video games




TÀI LIỆU SƯU TẬP
BỞI HCMUT-CNCP

Objective: Complete the game with the highest score



State: Raw pixel inputs of the game state

Action: Game controls e.g. Left, Right, Up, Down

Reward: Score increase/decrease at each time step



✓ Examples:
Go





Objective: Win the game!

State: Position of all pieces


Action: Where to put the next piece down

Reward: 1 if win at the end of the game, 0 otherwise

- ✓ Example:
- ❖ Put an agent in any room
 - ❖ Goal: go to Room 5 with fastest route



- ✓ State: Room 0, Room 1, ..., Room 5
- ✓ Action: Go to Room 0, Go to Room 1, ..., Go to Room 5
- ✓ Reward: matrix R



- ✓ Matrix Q: memory of what agent has learned through experience
 - ❖ Agent starts out knowing nothing
 - ❖ Q is initialized to zero



$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

- ✓ Defined:
 - ❖ States
 - ❖ Actions
 - ❖ Rewards matrix R
 - ❖ Matrix Q
- ✓ Training in progress
 - ❖ Updating matrix Q



✓ Utilize the Q matrix:

- ❖ Step 1: Set current state = initial state.
- ❖ Step 2: From current state, find the action with the highest Q value.
- ❖ Step 3: Perform action chosen in Step 2
- ❖ Step 4: Set current state = next state.
- ❖ Step 5: Repeat Steps 2, 3 and 4 until current state = goal state.



✓ Q learning algorithm:

Set the gamma parameter, and environment rewards in matrix R.

Initialize matrix Q to zero.

For each episode:

 Select a random initial state.

 While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

 End While

End For

discount factor

estimate of optimal future value

Q learning

✓ Q learning algorithm: gamma = 0.8, episode 1, initial state: 1

state = 1

action: go to 5

next_state = 5

For each episode:

Select a random initial state.

While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

End While

End For

$R =$

State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

$Q =$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

$Q =$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

Q learning

✓ Q learning algorithm: episode 2, initial state = 3

state = 3

action: go to 1

next_state = 1

For each episode:

Select a random initial state.

While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

End While

End For

0.8

100

$R =$

State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

$Q =$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

$Q =$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	80	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

Q learning

✓ Q learning algorithm: episode 2, initial state = 3

state = 1 action: go to 5

For each episode:

Select a random initial state.

While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

End While

End For

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 400 & 0 \\ 0 & 0 & 0 & 320 & 0 & 500 \\ 0 & 0 & 0 & 320 & 0 & 0 \\ 0 & 400 & 256 & 0 & 400 & 0 \\ 320 & 0 & 0 & 320 & 0 & 500 \\ 0 & 400 & 0 & 0 & 400 & 500 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

✓ Q learning algorithm:

Set the gamma parameter, and environment rewards in matrix R.

Initialize matrix Q to zero.

For each episode:

 Select a random initial state.

 While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

 End While

End For

discount factor

estimate of optimal future value

✓ Q learning algorithm:

Set the gamma parameter, and environment rewards in matrix R.

Initialize matrix Q to zero.

For each episode:

 Select a random initial state.

 While the goal state hasn't been reached.

- Select (randomly) one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute: $Q(\text{state}, \text{action}) \leftarrow (1 - \alpha) * Q(\text{state}, \text{action}) + \alpha * \{R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]\}$
- Set the next state as the current state.

 End While

End For

discount factor

estimate of optimal future value

✓ References

- ❖ http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture14.pdf
- ❖ <http://mnemstudio.org/path-finding-q-learning-tutorial.htm>

